

# Two Threads Beat One, and the Thread Hop Is a Wash Under Hawkes Arrivals: A Closed-Form and Empirical Case for Multi-Actor HFT Pipelines

Vincent Maciejewski\*

September 2026

## Abstract

The prevailing view in high-frequency-trading (HFT) system design is that the entire hot path — market-data decode, order-book application, strategy evaluation, order construction — should live on a single-threaded event loop, because every additional thread hop adds wall-clock latency that no benefit can recover. We show that this reasoning silently assumes a near-Poisson arrival process, and that it fails on real CME MDP3 market data, whose packet arrivals are strongly self-exciting (Hawkes-clustered) rather than Poisson.

Two consequences follow. First, on a clustered feed the end-to-end latency tail is dominated by cluster-induced queueing rather than by the per-hop cost, so a thread hop is not the first-order effect the single-thread argument assumes. Second, splitting the servicing chain into concurrent actor stages lowers the latency tail — because the queue that builds behind a burst is drained by several servers running in parallel — at the price of adding one hop to the median per stage; above a short-task floor set by the hop cost the tail reduction dominates and a multi-actor pipeline beats the single-thread chain on the tail.

The effect is governed by the arrival law, and the comparison against a matched-rate Poisson null separates two regimes that are easy to conflate. Over the band of per-message service times that an HFT hot path actually occupies — up to roughly  $20\mu\text{s}$  on this feed — the Poisson null has no tail at all: its  $p_{99}$  sits exactly at the service time, there is nothing for a pipeline to compress, and splitting is pure hop cost. The clustered feed has a substantial tail over that same band, so the entire win in the regime of interest is attributable to clustering rather than to adding threads. A Poisson process of the same average rate does eventually develop a tail — the paper should not be read as claiming otherwise — but only once ordinary utilisation queueing takes over, and for a matched-rate Poisson null that crossover is exactly located: by PASTA the  $q$ -quantile of latency equals the service time until the utilisation  $\rho = \bar{\lambda}T$  reaches  $1 - q$ , which the corpus confirms to three significant figures at  $q = 0.95, 0.99$  and  $0.999$ . At the corpus-median rate that puts the Poisson  $p_{99}$  crossover at  $T \approx 22\mu\text{s}$ , roughly three times the service time at which the clustered feed’s tail has already switched on, and past the crossover the clustered feed’s tail excess is still  $11\text{--}45\times$  the Poisson null’s. Clustering, not utilisation, is what a low-latency pipeline is queueing against. We give an analytic framework for the mechanism — a burst-limit throughput identity and an exact pathwise reduction of the deterministic-service tandem to a single bottleneck server — and confirm its predictions with simulations driven by a corpus of real CME MDP3 packet arrivals. We then cross-check the arrival-process result against an independent measurement taken on a live production receiver seven months after the corpus ends, on instruments the corpus does not contain and on a different arrival clock; the clustering statistics agree, and the production decoder’s measured service floor lands on the service time at

---

\*M2 Tech (16425640 Canada Inc.), Montreal. [majeski@gmail.com](mailto:majeski@gmail.com). KASPAR-HFT is open source at <https://github.com/vincent212/kaspar-hft>.

which our sweep says the latency tail switches on. Removing the constant-service assumption and re-running the sweep with service that depends on the number of messages a datagram carries leaves the design rule unchanged at every service time, and isolates a second tail mechanism — in-packet serialisation — that supplies the whole tail below the switch-on threshold and is the larger of the two at the production operating point, where it accounts for two thirds of the  $p_{99}$  excess. Companion code contains the message-tape parser, the per-window Hawkes fitter, and the tandem queueing simulator.

# Contents

|            |                                                                                     |           |
|------------|-------------------------------------------------------------------------------------|-----------|
| <b>1</b>   | <b>Introduction</b>                                                                 | <b>5</b>  |
| 1.1        | The design choice, and the safe default . . . . .                                   | 5         |
| 1.2        | Prior recommendations for single-thread hot paths . . . . .                         | 5         |
| 1.3        | What this paper does . . . . .                                                      | 7         |
| 1.4        | Contribution and prior-art positioning . . . . .                                    | 8         |
| 1.5        | Note on the use of AI in this paper . . . . .                                       | 8         |
| 1.6        | Notation . . . . .                                                                  | 9         |
| <b>I</b>   | <b>Analytical framework</b>                                                         | <b>10</b> |
| <b>2</b>   | <b>Analytical framework</b>                                                         | <b>10</b> |
| 2.1        | Burst-limit throughput theorem . . . . .                                            | 10        |
| 2.2        | Exact reduction of the tandem to one server, and the scaling bound . . . . .        | 12        |
| 2.3        | What the theory says, in plain terms . . . . .                                      | 15        |
| <b>II</b>  | <b>Empirical experiment</b>                                                         | <b>18</b> |
| <b>3</b>   | <b>Experiment setup</b>                                                             | <b>18</b> |
| 3.1        | Data corpus . . . . .                                                               | 19        |
| 3.2        | Packet arrivals are Hawkes-clustered (empirical prerequisite) . . . . .             | 20        |
| 3.3        | Simulator . . . . .                                                                 | 21        |
| 3.4        | Pipelining is what makes multi-stage work . . . . .                                 | 21        |
| 3.5        | Sweep design . . . . .                                                              | 22        |
| <b>4</b>   | <b>Empirical results</b>                                                            | <b>22</b> |
| 4.1        | Tail collapse under Hawkes, and its near-absence under Poisson at HFT service times | 22        |
| 4.2        | Median cost is linear in $N$ , exactly as $(N - 1)h$ . . . . .                      | 25        |
| 4.3        | Design equation . . . . .                                                           | 25        |
| 4.4        | Across windows . . . . .                                                            | 27        |
| 4.5        | Where the tail switches on . . . . .                                                | 28        |
| 4.6        | Conditioning on intensity and criticality . . . . .                                 | 28        |
| 4.7        | Sensitivity to $h$ . . . . .                                                        | 31        |
| <b>5</b>   | <b>Evaluation: do empirics match the analytical predictions?</b>                    | <b>31</b> |
| <b>6</b>   | <b>Cross-validation: the same arrival process, measured live</b>                    | <b>33</b> |
| 6.1        | Two measurements, independent where it counts . . . . .                             | 33        |
| 6.2        | What agrees . . . . .                                                               | 33        |
| 6.3        | The service floor lands on the switch-on threshold . . . . .                        | 34        |
| 6.4        | Where the two disagree, and what each correction is . . . . .                       | 35        |
| 6.5        | What the live data contradicts, and what survives . . . . .                         | 36        |
| 6.6        | Span-aware service: re-running the sweep with the assumption removed . . . . .      | 37        |
| <b>III</b> | <b>Best practices: implications for HFT systems design</b>                          | <b>39</b> |

|           |                                                                       |           |
|-----------|-----------------------------------------------------------------------|-----------|
| <b>7</b>  | <b>The reference instance: KASPAR-HFT and its fast_send pipeline</b>  | <b>40</b> |
| <b>8</b>  | <b>Implementation on KASPAR-HFT: two dispatch paths, one handler</b>  | <b>41</b> |
| 8.1       | Receiver transparency . . . . .                                       | 42        |
| 8.2       | Why measurement, not local reasoning, must decide . . . . .           | 42        |
| 8.3       | Operational discipline . . . . .                                      | 43        |
| 8.4       | Separation of code construction from latency optimization . . . . .   | 43        |
| 8.5       | Contrast with non-fungible frameworks . . . . .                       | 44        |
| <b>9</b>  | <b>Other architectural considerations to re-evaluate under Hawkes</b> | <b>44</b> |
| <b>10</b> | <b>Recommendations for HFT system designers</b>                       | <b>46</b> |
| 10.1      | What the paper leaves untouched . . . . .                             | 46        |
| 10.2      | What the paper does argue against . . . . .                           | 46        |
| 10.3      | Recommended design flow . . . . .                                     | 47        |
| <b>IV</b> | <b>Conclusions</b>                                                    | <b>49</b> |
| <b>11</b> | <b>Conclusion</b>                                                     | <b>49</b> |
| <b>12</b> | <b>Future work</b>                                                    | <b>51</b> |
| <b>13</b> | <b>Reproducibility</b>                                                | <b>51</b> |
|           | <b>Appendices</b>                                                     | <b>56</b> |
| <b>A</b>  | <b>The Borel batch-size tail at fixed branching ratio</b>             | <b>56</b> |
| <b>B</b>  | <b>Numerical check of Theorem 1</b>                                   | <b>56</b> |
| <b>C</b>  | <b>Numerical check of the unequal-stage claims</b>                    | <b>57</b> |

# 1 Introduction

## 1.1 The design choice, and the safe default

The design of the ingress-to-order pipeline in a high-frequency trading system is a choice between two archetypes: the single-threaded event loop that processes every step of every message on one core (SBE decode, book application, strategy evaluation, order construction, all under one call chain, taking a total service time we denote  $T$ ), and the multi-actor pipeline — a *tandem* of  $N$  servers in series — in which those steps are split across  $N$  separate actors, each on its own thread with its own mailbox, performing its part of the work and handing the message to the next across one of  $N - 1$  explicit hops whose per-hop wall-clock cost we denote  $h$ . It is not really a debate: almost everyone in the low-latency HFT community is in the single-thread camp, and single-threaded is generally viewed as the safe choice. The argument for it cites the cost of a mailbox hop — typically a mailbox enqueue plus a cache-line transfer between the sender’s core and the receiver’s core, plus the cost of the receiver being scheduled or spinning — and holds that  $(N - 1)h$  adds directly to every message’s end-to-end latency, so every extra stage in the pipeline is a fixed tax that no throughput or tail benefit can overcome.

That argument implicitly assumes an arrival process for which single-server throughput  $1/T$  is comfortably larger than the peak arrival rate. On CME MDP3 the arrival process is strongly self-exciting: packets arrive in Hawkes-clustered bursts [24–26]. During such bursts the instantaneous arrival rate transiently exceeds the single-thread event loop’s service rate  $1/T$ , and cluster-induced queueing accumulates a tail that dwarfs  $h$  by one to two orders of magnitude. This paper’s central empirical observation, on 276 sessions (3512 thirty-minute windows) of NQ front-month packet arrivals, is that this cluster tail is what dominates observed end-to-end latency: the 99th-percentile latency ( $p_{99}$ ) of a constant-service queue driven by real packet arrivals is 2.2 times the service floor at the corpus-median window for a  $16\mu\text{s}$  task and 4.6 times for a  $32\mu\text{s}$  task, and 2.5 and 6.1 times at the 95th-percentile window. Under a rate-matched Poisson null of the same arrival count the  $p_{99}$  still equals the median at  $16\mu\text{s}$  and is only 1.3 times it at 32, against the Hawkes arm’s 4.6; over this band the tail is a Hawkes-cluster effect almost in its entirety.

Given that the tail is cluster-driven, the natural question is whether splitting the pipeline into  $N > 1$  stages — each running as an independent actor on its own core, with per-stage service  $T/N$ , so that the tail-driving cluster-wait is now split across  $N$  smaller queues — reduces the tail. If it does, and if the additional median cost  $(N - 1)h$  is small relative to the tail savings, the multi-actor architecture is the better choice for any operator who weights tail latency over median latency.

## 1.2 Prior recommendations for single-thread hot paths

The single-thread rule has a specific lineage in the low-latency-systems community, and we name that lineage here because the paper is arguing against it directly. In reverse chronological order, the lineage has three layers: the HFT-specific codification of the last decade, the general low-latency-systems and event-driven-server culture from which HFT inherited it, and one prominent published counter-argument that the field discarded.

**Modern HFT-community codification (2010–present).** The founding statement of the current HFT-community formulation is Thompson’s Single Writer Principle [34] — “a single-writer design consistently outperforms a ‘properly’ lock-free or wait-free multi-writer design on most realistic workloads” — motivated by and applied in the LMAX Exchange architecture [35, 36]. Rigtorp’s widely cited low-latency tuning guide [37] formalises the operational rule (“run a single thread in

SCHED\_OTHER per core and use busy waiting/polling”), and Chronicle Software’s public engineering materials [38] adopt the same architecture for market-data connectors (“within a single event loop . . . busy-spinning and affinity locks”), pinning Chronicle Queue’s design to a single writer per queue. The Aeron messaging library [39] inherits the same lineage. The academic side is thinner, but Kwan et al. [40] codify the pattern as the standard low-latency-HFT design pattern, and recent practitioner posts [41] restate it explicitly (“all critical trading logic runs in a single-threaded event loop, eliminating context switches and lock contention entirely”). Inside HFT specifically, the same recommendation has been reinforced by two decades of private prop-trading arms-race practice among the Chicago desks (Getco, Jump, Optiver, DRW; late 2000s), reported second-hand in industry retrospectives but without an open-literature reference.

**Older event-driven-server lineage (1995–2010).** HFT did not invent this rule; it inherited it from the event-driven-server culture. The single-threaded event loop as an architectural pattern dates at least to Schmidt’s Reactor pattern [42] in the mid-1990s. It was sharpened by Kegel’s C10K essay [43], which named the OS thread-per-connection model as the bottleneck of scalable network servers and argued for event-driven designs. The pattern went mainstream in production with Sysoev’s nginx web server [44], whose single-threaded worker process is explicitly justified as context-switch avoidance, and later with Dahl’s Node.js runtime [45], which pushed the single-threaded event loop into ordinary application code.

**Pro-pipelining minority — outside HFT.** A legitimate multi-decade pro-pipelining tradition exists in the datacenter and networking-systems literature; it has been the field’s non-single-thread minority throughout the period covered by the previous two paragraphs. The earliest prominent published argument is the Staged Event-Driven Architecture (SEDA) of Welsh et al. [46], which advocated exactly the multi-stage pipeline of independently scheduled event handlers that the present paper’s tandem design revives. Welsh’s 2010 retrospective [47] is often cited as a walk-back to single-thread, but it is actually a hedged pro-pipelining recommendation: it advocates decoupling stages from thread pools, grouping multiple stages within a single thread-pool domain where latency is critical, and putting a separate thread pool and queue in front of any stage with long or nondeterministic runtime — effectively the present paper’s design rule in different vocabulary. Beyond SEDA the pro-pipelining tradition includes the Click Modular Router [48] for staged packet processing; Naiad [49] for low-latency dataflow; Snap [50] for Google’s multi-engine network stack; and the datacenter microsecond-tail-latency scheduler line — Shinjuku [51], Shenango [52], Caladan [53], Perséphone [54] — which argues for centralised multi-core dispatch specifically to control the microsecond-scale tail. Perséphone in particular is targeted at “wide service-time distributions,” the closest existing framing to our arrival-law claim; but the datacenter line frames the problem as heavy-tailed *service*, not clustered *arrivals*. DPDK explicitly supports both run-to-completion and pipeline modes and its documentation recommends pipeline mode “if some set of packets require longer to process” [55]. Inside HFT specifically the pro-pipelining position is rare in the open literature; an obscure Chinese patent [56] describing a pipelined FAST market-data decoder is a lonely commercial artefact of it.

None of the pro-pipelining sources above ties the design recommendation to Hawkes / self-exciting / near-critical arrival dynamics. The closest theoretical hooks are the MMPP-input tandem work of Kim and Chydziński [16] already cited in Section 1.4 and the batch-arrival tandem phase-transition analysis of Vinnikov [57]. The present paper’s arrival-law-specific mechanism — that the recommendation flips when arrivals are near-critically self-exciting — remains uncontested by prior work.

Yours truly gets a footnote on that list, but with an important caveat.<sup>1</sup> An earlier paper by the present author [32] did not argue for the single-thread principle in its own right. That paper is defensive rather than prescriptive: the actor model had long been dismissed as unsuitable for HFT because of asynchronous-dispatch overhead, and `fast_send` is presented there as the mechanism that eliminates the overhead for co-located actors and thereby makes the actor model usable *inside* a servicing chain that the prevailing view of the time had already placed on a single thread. The reference architecture in that paper — reader  $\rightarrow$  buffer via asynchronous `send`, decode/book/signal chain via `fast_send` on one thread, signal  $\rightarrow$  order-out via asynchronous `send` — takes the collapsed single-thread servicing chain as the widely accepted design context of the time and delivers a way to compose it out of actors at zero context-switch cost. It does not compare single-thread against more-asynchronous topologies and pick single-thread; it takes the topology as given. The present paper is a retraction of that design context, not of the `fast_send` mechanism: on Hawkes-clustered feeds, the aggregate servicing chain should be split into asynchronous stages once its total service time becomes long compared to the Hawkes-cluster tail. `fast_send` remains the right mechanism for chaining actors that have been placed together on one thread; what changes is where the thread boundaries inside the servicing chain should sit. The corrected design rule is stated in Section 11.

### 1.3 What this paper does

We answer that question along three axes.

- **Theory.** Two closed-form results (Section 2) predict the pipeline behaviour. Proposition 1 shows that under any arrival process an  $N$ -stage tandem enjoys an  $N$ -factor throughput speedup on isolated bursts. Theorem 1 shows that a deterministic-service tandem reduces exactly to a single server with the largest stage’s service, and that the end-to-end wait is bounded pathwise by  $1/N$  of the single-stage wait plus  $(N - 1)h$ , at every quantile and under every arrival sequence. Corollary 2 shows that under a rate-matched Poisson null the single-stage wait at the  $q$ -quantile is already zero whenever the utilisation satisfies  $\rho < 1 - q$  — which at this feed’s rate covers  $p_{99}$  for every task up to  $16\mu\text{s}$  — so over that band there is nothing to compress and splitting costs exactly  $(N - 1)h$ . The Hawkes–Oakes cluster representation then explains why the single-stage wait on a real feed is large — one to two orders of magnitude above the per-hop cost  $h$  that the single-thread argument worries about. A self-exciting feed is equivalent to a Poisson stream of background “parent” arrivals, each of which sets off a cluster of follow-on packets — a parent is followed by children, those by grandchildren, and so on. How large these clusters get is controlled by the branching ratio  $n$ , the mean number of direct children per packet: the mean cluster holds  $1/(1 - n)$  packets and the largest clusters routinely reach a scale of  $(1 - n)^{-2}$  packets (the cluster-size law is the Borel distribution; Appendix A). At the fitted  $n \approx 0.8$  that is a mean of about five and a scale of about twenty-five packets, arriving essentially back-to-back. Such a burst of  $K$  packets leaves its last member waiting behind  $K - 1$  predecessors on a single server (Proposition 1), so the single-stage wait grows with cluster size — the tail the tandem then compresses.
- **Simulation.** We build an  $N$ -stage tandem Lindley recursion on real CME MDP3 packet-arrival streams at a hop cost  $h = 1.7\mu\text{s}$  (the one-way async mailbox cost derived by halving

---

<sup>1</sup>A related author’s-aside: the tail compression reported by this paper had in fact been observed empirically in production well before its mechanism was understood. Splitting a long task into shorter asynchronously connected stages was known to help the tail on a per-strategy basis in the author’s own operational experience, and the working assumption at the time was that the single-thread camp must simply be writing faster inner-loop code — not that the architectural recommendation was itself arrival-law-dependent. The present paper is the retrospective explanation of that anecdote.

the round-trip figure in Table 4 of Maciejewski [32]) and a service floor swept across  $T \in \{2, 4, 8, 16, 32, 64, 128\} \mu s$  to cover the regime from tasks below the hop cost (where no split helps) through tasks well above it (where the tail-win compounds). We sweep  $N \in \{1, 2, 4, 8\}$  per scenario on each of the 3512 thirty-minute RTH windows of the corpus. Each window is also simulated under a uniform-shuffle Poisson null of the same arrival count as an empirical counterfactual (Section 3.5).

- **Design equation for practitioners.** From the empirical result we extract a concrete set of design rules for a latency-critical servicing chain (Sections 7 and 8): split a stage only when its service time exceeds the feed’s tight-gap floor, about  $7.5 \mu s$  on this corpus, since below it there is no queueing tail to compress; choose the partition with the smallest achievable largest stage  $s_{\max}$ , since the tail depends only on  $s_{\max}$  and a cut that does not lower it merely buys a hop; put that largest stage first, at ingress; keep sub-tasks below the bottleneck together on one thread via `fast_send`, using an asynchronous `send` only at boundaries that lower  $s_{\max}$ ; and settle each boundary by end-to-end measurement against a representative feed rather than by local reasoning at a single call site.

## 1.4 Contribution and prior-art positioning

A literature search covering foundational tandem theory (Jackson, Kelly, Burke, Kleinrock, Whitt), tandem theory under bursty arrivals (Foss–Korshunov, Baccelli–Foss–Lelarge, Heindl, Kim–Chydzinski, Dębicki–Mandjes), Hawkes single-queue and adjacent (Daw–Pender, Chen–Blanchet, Koops–Boxma–Mandjes, Sheldon et al., Gao–Zhu, Bo Li, and 2023–2026 follow-ons), and empirical/applied HFT and RPC-pipeline work (ABIDES, ABIDES-MARL, JAX-LOB, Aquilina–Budish–O’Neill, Dean–Barroso, Parsimon, and the datacenter tail-latency crowd) confirmed that there is no direct prior art on either of the paper’s two claims. The first is theoretical: that the choice between a single-queue (single-thread) architecture and an  $N$ -stage tandem (multi-actor pipeline) is not fixed but *reverses with the arrival law*. At a given service time — and we mean the ones an HFT hot path actually occupies, single-digit to low-tens of microseconds — Poisson arrivals make the single queue the better choice, because at that utilisation there is no wait to compress and each additional stage only adds hop latency, while Hawkes-clustered arrivals at the *same mean rate* make the tandem win on the tail, because the cluster wait it compresses outweighs the hops it adds. The reversal is between arrival laws at fixed rate and service time, not a claim that Poisson queues never build tails: push the service time far enough and a Poisson queue tails too, just an order of magnitude later and far more mildly. We have found no prior work stating that the architectural recommendation itself flips sign between these two arrival laws. The second is empirical: a tandem-queue latency sweep — over a range of service floors and stage counts — driven by real CME MDP3 packet timestamps, controlled against a rate-matched Poisson null, and calibrated with a hop cost measured on a production HFT actor framework. We have found no prior study that combines a real exchange-calibrated tandem sweep with both that Poisson counterfactual and a measured framework hop cost. The full reference list appears in Section 13.

## 1.5 Note on the use of AI in this paper

This paper was produced with Anthropic’s Claude (via Claude Code) used under the author’s direction for code generation — the message-tape parser, the packet-Hawkes fitter, and the tandem-Lindley simulator run on the 3512-window corpus — and for exploratory analysis, mathematical drafting, prior-art literature search, prose editing, and figure preparation. The analytical results of Part I, including the proof of Theorem 1, were drafted with the model and then independently

checked by the author. The author independently reviewed all mathematical arguments, experimental methodology, numerical results, and citations, and takes sole responsibility for the contents of the paper.

## 1.6 Notation

The following symbols are used throughout.

| Symbol                       | Meaning                                                                                                              |
|------------------------------|----------------------------------------------------------------------------------------------------------------------|
| $T$                          | total service time per message on a single-stage pipeline, in $\mu\text{s}$                                          |
| $N$                          | number of stages in the tandem pipeline (one actor per stage)                                                        |
| $T/N$                        | per-stage deterministic service time in the $N$ -stage tandem                                                        |
| $h$                          | per-hop wall-clock cost of the async mailbox handoff between adjacent stages                                         |
| $k$                          | size of an arrival burst (number of packets in one cluster)                                                          |
| $\tau_k^{(N)}$               | exit time of the last message of a $k$ -burst under $N$ stages                                                       |
| $s_i, s_{\max}$              | per-stage services of an unequal tandem, $\sum_i s_i = T$ , and their maximum                                        |
| $w_j(s)$                     | Lindley waiting time of message $j$ at a single deterministic server with service $s$                                |
| $W_j^{(N)}, W_q^{(N)}$       | end-to-end wait of message $j$ (latency minus $T$ , hops included) under $N$ stages, and its $q$ -quantile           |
| $p_q$                        | $q$ -th percentile of the end-to-end latency distribution ( $p_{50}$ , $p_{99}$ , etc.)                              |
| $\Delta(N)$                  | clustering tail excess $p_{99}(N) - p_{50}(N)$ (§4.3)                                                                |
| $\Delta_{\text{single}}(s)$  | single-server clustering tail excess at deterministic service $s$ ; $\Delta(N) = \Delta_{\text{single}}(T/N)$ (§4.3) |
| $\gamma$                     | fitted exponent in $\Delta(N) \approx \Delta(1)N^{-\gamma}$                                                          |
| $\kappa$                     | operator's tail-versus-median weight in the design equation                                                          |
| $N^*$                        | stage count minimising the design objective $(N-1)h + \kappa \Delta(N)$ (§4.3)                                       |
| $\mu, \alpha, \beta$         | Hawkes background rate, jump size, and decay rate (§3.2)                                                             |
| $n = \alpha/\beta$           | Hawkes branching ratio                                                                                               |
| $K$                          | Borel-distributed offspring cluster size (§2.2)                                                                      |
| $\bar{\lambda}_{\text{pkt}}$ | mean packet-arrival intensity per 30-min window, packets per second                                                  |
| $n_{\text{pkt}}$             | branching ratio of the packet-arrival Hawkes fit per window                                                          |
| $n_{\text{pkt}}^{\#}$        | per-window packet count                                                                                              |
| $\rho$                       | utilisation, $\rho = \bar{\lambda}_{\text{pkt}} \cdot T$                                                             |
| $\tau$                       | bin width of the Fano factor, in seconds (distinct from the service time $T$ )                                       |
| $F(\tau)$                    | Fano factor at bin width $\tau$ : $\text{Var}[C_\tau]/\text{E}[C_\tau]$                                              |
| $H$                          | Hurst exponent, from log-log slope of $F(\tau)$                                                                      |
| $\sigma_i, \bar{\sigma}$     | span of packet $i$ (SBE messages per datagram) and its window mean (§6.6)                                            |
| $S_i$                        | span-dependent service time of packet $i$ (§6.6)                                                                     |

The per-window packet count and the packet-Hawkes branching ratio would both naturally be written  $n_{\text{pkt}}$ ; the paper reserves  $n_{\text{pkt}}$  for the branching ratio and writes  $n_{\text{pkt}}^{\#}$  for the count.

## Part I

# Analytical framework

## 2 Analytical framework

Throughout, the *tandem* is the paper’s model of the multi-actor pipeline: the total per-message service time  $T$  is split into  $N$  stages placed in series, each stage an independent deterministic-service server running on its own thread that performs its share of the work and passes the message to the next stage across a mailbox hop of cost  $h$ . The single-thread event loop is the  $N = 1$  case, in which one server performs all of  $T$ ; adding stages shortens each server’s service but inserts a hop between neighbours.

A baseline proposition, a theorem, and a corollary anchor the paper’s simulation. Proposition 1 is the isolated-burst identity. A message’s end-to-end *latency* is its *service time* — the time to actually run it,  $T$  — plus its *queueing wait*, the time it sits behind predecessors before service begins. Running the  $N$  stages concurrently raises throughput from  $1/T$  to  $N/T$ , so an  $N$ -stage tandem clears a back-to-back burst  $N$  times faster: the service time stays  $T$ , while the last message’s queueing wait — and hence its latency, once the wait dominates — shrinks by a factor of  $N$ . This is the classical pipeline speedup, taken here as an obvious baseline without proof.

For example, split a service- $T$  server into two equal  $T/2$  stages and feed it a burst of  $k$  back-to-back messages. All of the queueing now happens at the first stage: its last message waits behind  $k - 1$  predecessors that each occupy it for only  $T/2$ , a wait of  $(k - 1)T/2$  against  $(k - 1)T$  on the single server. The second stage never queues — stage-one departures arrive spaced  $T/2$  apart, exactly its service rate, so each message passes straight through for a further  $T/2 + h$ . The wait is not spread across two smaller queues; it stays in one bottleneck stage that is now twice as fast. The halving depends on the split being equal: cutting  $T = 10$  as  $7 + 3$  leaves a bottleneck of 7, so the wait drops only to  $7/10$ , and in general the tail scales with the largest stage  $s_{\max}$  rather than with the number of stages — the point Theorem 1 makes exact. Theorem 1 is the general statement: the tandem’s end-to-end wait equals the wait at a single server with the largest stage’s service plus  $(N - 1)h$ , exactly and on every sample path, and that single-server wait is at most  $1/N$  of the wait at service  $T$ . Both are elementary and arrival-law independent; Theorem 1 rests on a reduction theorem of Friedman [11] and on the monotonicity of the Lindley recursion. Corollary 2 is the Poisson null: it locates exactly where the Poisson wait switches on, at  $\rho = 1 - q$ , and at the utilisation of a CME packet feed that puts the single-stage  $p_{99}$  wait at zero for every task up to  $16\mu\text{s}$ , so over that band splitting costs exactly  $(N - 1)h$  and buys nothing. The arrival law enters only through the size of the single-stage wait that the  $1/N$  is applied to, and the Hawkes–Oakes cluster representation explains why that wait is large on a real feed. The literature search summarised in Section 1.4 found no published statement of the architecture recommendation flipping with the arrival law — single thread under Poisson, multi-actor tandem under Hawkes-clustered arrivals; the results below are the paper’s formal content, and Part II measures the slack in the bound on real arrivals.

### 2.1 Burst-limit throughput theorem

**Proposition 1** (Burst-limit  $N$ -factor speedup). *Consider an isolated cluster of  $k$  arrivals whose  $j$ -th member arrives no later than  $(j - 1)T/N$  after the first, so that no server idles during the burst (an instantaneous batch is the canonical case). Let  $\tau_k^{(1)}$  denote the exit time of the last arrival under a single-stage deterministic-service queue with service  $T$ , and let  $\tau_k^{(N)}$  denote the exit time under an*

$N$ -stage tandem with per-stage deterministic service  $T/N$  and per-stage hop delay  $h$ . Then

$$\tau_k^{(1)} = kT, \quad \tau_k^{(N)} = T + (k-1) \frac{T}{N} + (N-1)h, \quad (1)$$

and in particular

$$\lim_{k \rightarrow \infty} \frac{\tau_k^{(1)}}{\tau_k^{(N)}} = N. \quad (2)$$

**Remark.** Proposition 1 is the classical pipelined-tandem throughput result [8, 9]; we state it without proof and use it only to fix the baseline against which the arrival-dependent tail reduction of Theorem 1 is measured. What is paper-specific is the numerical calibration of  $T$  and  $h$  on a production HFT actor framework and the observation that under real CME MDP3 arrivals the ratio  $\tau_k^{(1)}/\tau_k^{(N)}$  is what governs the observed tail-latency compression.

**Reading Equation (1).** In the paper’s HFT setting we take  $T = 10\mu\text{s}$  as a typical HFT task on a single-thread event loop, and  $k$  is the size of a Hawkes-cluster burst that hits the pipeline back-to-back. The single-stage formula  $\tau_k^{(1)} = kT$  says that in an event-loop design the last message of a  $k$ -burst waits behind  $k-1$  predecessors, each of which spends the full  $T$  on the one server; a 10-packet cluster on a  $10\mu\text{s}$  server therefore has a last-message latency of  $100\mu\text{s}$ . The tandem formula  $\tau_k^{(N)} = T + (k-1)(T/N) + (N-1)h$  decomposes the last-message latency into three physical pieces: (i)  $T$ , the total service work every message must traverse regardless of  $N$ ; (ii)  $(k-1)(T/N)$ , the pipelined wait behind  $k-1$  predecessors at throughput  $N/T$  rather than  $1/T$ ; and (iii)  $(N-1)h$ , the one-time hop-cost tax paid to traverse  $N-1$  mailbox boundaries on the way through. The single message ( $k=1$ ) pays only (i) and (iii), never (ii). The whole point of the paper is that under Hawkes arrivals term (ii) is what blows up on a single stage and shrinks by a factor of  $N$  under a tandem, while term (iii) is only a modest additive tax when the framework’s  $h$  is small.

**Reading Equation (2).** The ratio  $\tau_k^{(1)}/\tau_k^{(N)}$  tends to  $N$  as the cluster size  $k$  grows, meaning that in the burst-saturation limit a tandem of  $N$  actors clears the burst  $N$  times faster than an event loop. This is the classical pipeline throughput speedup restated for HFT: a 4-actor pipeline drains a large burst in one-quarter of the wall-clock time that a single-thread event loop would take, and a 12-actor pipeline in one-twelfth. It is worth emphasising that this speedup is achieved with the same total service work per message: the pipeline does not reduce how much decode/book/strategy work happens per event, it only overlaps the work across cores so the queue behind a big cluster clears in parallel. The arrival law does not enter Equation (2); what the arrival law controls is *how often* the burst regime is entered, and Theorem 1 shows the same factor governs every message on every sample path.

**On what  $T$  is, and on the natural stage boundaries.** The paper adopts  $T = 10\mu\text{s}$  as a ballpark typical HFT single-thread event loop. Any real HFT pipeline of that duration is already a sequential chain of internally atomic sub-tasks (SBE decode, order-book application, strategy touch, order construction on CME MDP3, per Section 8), so the  $N$  stages of the tandem are placed at the existing boundaries between those sub-tasks; the paper is not subdividing any single sub-task, and in particular is not subdividing the SBE decode, which is a linear pass over the wire bytes of a single message and cannot be split. Ten microseconds is a round-number ballpark chosen for exposition; none of the paper’s qualitative results depend on it. The relevant knob for the design equation is the ratio  $T/h$ : doubling  $T$  at fixed  $h$  doubles the cluster-wait terms of Equation (1) and leaves the

hop tax unchanged. The floor below which splitting cannot pay (Section 4.3) turns out to be set by the feed rather than by  $h$ : it is the arrival stream’s tight-gap floor, and on this corpus it binds well above  $2h$ . The paper writes “per-stage service  $T/N$ ” as a simplifying modelling convention that assumes equal-sized stages; in reality the per-stage services  $s_i$  sum to  $T$  but are unequal, and the achievable pipeline throughput is  $1/\max_i s_i$  rather than  $N/T$ . Theorem 1 handles the unequal case exactly: the factor becomes  $s_{\max}/T$  in place of  $1/N$ , the tail still contracts under Hawkes, and the median still grows linearly at slope  $h$  per hop.

This is the mechanism. It applies to any deterministic-service tandem regardless of the arrival law; the arrival law determines *how often* large- $k$  clusters occur, and the next subsection makes the per-burst factor a pathwise bound.

## 2.2 Exact reduction of the tandem to one server, and the scaling bound

Proposition 1 treats one isolated burst. The result that carries the paper’s design claim is that the same  $N$ -factor holds pathwise for every message under every arrival sequence, at the level of the end-to-end wait. Fix an arrival sequence  $a_1 \leq a_2 \leq \dots$  (any point process, clustered or otherwise, with batches allowed as coincident times) and write  $\delta_j = a_{j+1} - a_j \geq 0$ . For a single FIFO server with deterministic service  $s$  fed by this sequence, the waiting time of message  $j$  obeys the Lindley recursion [12]

$$w_1(s) = 0, \quad w_{j+1}(s) = (w_j(s) + s - \delta_j)^+. \quad (3)$$

Define the end-to-end wait of message  $j$  in the  $N$ -stage tandem as  $W_j^{(N)} = (\text{exit time}) - a_j - T$ , i.e. latency in excess of the total service work; under  $N = 1$  this is  $W_j^{(1)} = w_j(T)$ .

**Theorem 1** (Reduction and scaling bound). *Consider an  $N$ -stage tandem of FIFO servers with deterministic per-stage services  $s_1, \dots, s_N$  summing to  $T$ , deterministic hop delay  $h$  between consecutive stages, unbounded buffers, and an arbitrary arrival sequence. Let  $s_{\max} = \max_i s_i$ . Then for every message  $j$  and every sample path,*

$$W_j^{(N)} = w_j(s_{\max}) + (N - 1)h, \quad (4)$$

$$w_j(s') \leq \frac{s'}{s} w_j(s) \quad \text{for all } 0 < s' \leq s, \quad (5)$$

and consequently

$$W_j^{(N)} \leq \frac{s_{\max}}{T} W_j^{(1)} + (N - 1)h = \frac{W_j^{(1)}}{N} + (N - 1)h \quad \text{under equal stages } s_i = T/N. \quad (6)$$

Since Equation (6) holds pathwise, it holds for every quantile of the stationary distribution whenever one exists:  $W_q^{(N)} \leq W_q^{(1)}/N + (N - 1)h$  for all  $q$ . The same inequality holds for the empirical (sample) quantiles computed over any finite trace, directly from the pathwise bound and without invoking stationarity; the sample-quantile form is what Part II measures, and it should be read as a statement about the observed windows rather than about a stationary-distribution quantile or a finite-sample maximum.

*Proof.* Equation (4), equal stages. Departures from a FIFO server with deterministic service  $s$  are spaced at least  $s$  apart, because each departure requires  $s$  of uninterrupted service after the previous one. A FIFO server with deterministic service  $s$  whose arrivals are spaced at least  $s$  apart never queues: by induction, message  $j$  departs at its arrival time plus  $s$ , which is at or before the arrival of message  $j + 1$ . Stage 1 is a single server with service  $T/N$  fed by the original sequence, so message  $j$

leaves it at  $a_j + w_j(T/N) + T/N$ . The arrival sequence at stage 2 is the stage-1 departure sequence shifted by  $h$ , spaced at least  $T/N$  apart, so stage 2 adds exactly  $h + T/N$  and its departures are again spaced at least  $T/N$  apart; the same holds at each subsequent stage. Summing, the exit time is  $a_j + w_j(T/N) + N \cdot T/N + (N-1)h$ , and subtracting  $a_j + T$  gives Equation (4) with  $s_{\max} = T/N$ .

Equation (4), unequal stages. By the order-invariance theorem of Friedman [11], for a tandem of FIFO stages with constant service times, unlimited queues, and any arrival sequence, each customer's time in system is independent of the order of the stages. Deterministic hop delays preserve this: a hop is a constant shift of one stage's departure sequence into the next stage's arrival sequence, and constant shifts commute with the Lindley recursion, so reordering stages together with their hops leaves every exit time unchanged. Reorder so that the stage with service  $s_{\max}$  comes first. Its departures are spaced at least  $s_{\max}$  apart, and every subsequent stage has service  $s_i \leq s_{\max}$ , so by the no-queueing argument above each of them adds exactly  $s_i + h$  to every message. The exit time is  $a_j + w_j(s_{\max}) + \sum_i s_i + (N-1)h$ , and subtracting  $a_j + T$  gives Equation (4).

Equation (5). By induction on  $j$  with  $w_1(s') = 0 = (s'/s)w_1(s)$ . The map  $x \mapsto (x + s' - \delta_j)^+$  is nondecreasing, so

$$w_{j+1}(s') = (w_j(s') + s' - \delta_j)^+ \leq \left(\frac{s'}{s}w_j(s) + s' - \delta_j\right)^+ \leq \left(\frac{s'}{s}w_j(s) + s' - \frac{s'}{s}\delta_j\right)^+ = \frac{s'}{s}w_{j+1}(s),$$

where the second inequality uses  $\delta_j \geq (s'/s)\delta_j$ , which holds because  $\delta_j \geq 0$  and  $s' \leq s$ . Combining Equations (4) and (5) with  $s' = s_{\max}$ ,  $s = T$  gives Equation (6).  $\square$

**Remark** (What the theorem says and what it does not). Equation (4) says all queueing in an equal-stage deterministic tandem happens at stage 1; stages  $2, \dots, N$  contribute exactly  $T/N + h$  each to every message. Proposition 1 is the isolated-burst special case, where Equation (6) holds with equality. Equation (6) is a bound rather than an identity because the tandem runs its bottleneck stage at utilisation  $\rho/N$  while the single stage runs at  $\rho$ : the induction step in the proof of Equation (5) gives away  $(1 - 1/N)\delta_j$  at every message  $j$  that the single server is still serving when message  $j+1$  arrives, so the tandem clears each inter-arrival gap  $N$  times more effectively than the scaled bound credits it for. The paper's empirical  $\gamma$  (Section 4.3), defined by  $W_q^{(N)} - (N-1)h \approx N^{-\gamma}W_q^{(1)}$ , therefore satisfies  $\gamma \geq 1$  in the model, and the excess  $\gamma - 1$  measures that slack. Nothing in the theorem depends on the arrival law; what the arrival law determines is the size of  $W_q^{(1)}$  that the factor  $1/N$  is applied to, which is the subject of the next remark. The hop cost  $h$  enters every quantile as exactly  $(N-1)h$ , in the median and in the tail alike.

**Remark** (Unequal stages: only the largest stage matters). Real cut points do not give equal stages, and Equation (4) says what that costs: the end-to-end wait of any deterministic tandem equals the wait at a single server with service  $s_{\max}$ , plus  $h$  per hop. Three consequences follow, none of which the equal-stage form makes visible.

1. The tail factor is  $s_{\max}/T$ , and the stage count enters only through the hop tax. Splitting a  $10\mu\text{s}$  task as  $7 + 3$  compresses the tail excess to at most 70% of single-stage, against 50% for  $5 + 5$ , and both pay one  $h$ .
2. The order of the stages is irrelevant to end-to-end latency:  $7 + 3$  and  $3 + 7$  give identical latencies on every sample path. Order determines only where the queue physically sits (below).
3. A cut that does not lower  $s_{\max}$  buys nothing. Refining  $7 + 3$  into  $7 + 2 + 1$  adds a hop and leaves every quantile of the wait unchanged, so end-to-end latency rises by exactly  $h$ . Conversely, sub-tasks can be merged onto one thread (via `fast_send`) without any tail cost as long as the

merged service does not exceed  $s_{\max}$ . The design problem is therefore to choose the partition with the smallest  $s_{\max}$  reachable at the natural cut points, and then the fewest hops that realise it.

Appendix C checks all three numerically.

**Remark** (Where the burst goes: the bottleneck absorbs the clustering). The proof of Equation (4) rests on one fact: departures from a deterministic server with service  $s$  are spaced at least  $s$  apart. Applied to a Hawkes input, that fact says what the bottleneck stage does to the clustering. Upstream of the bottleneck the stream is the raw feed, with intra-cluster inter-arrival times far below  $s_{\max}$  and instantaneous intensity far above  $1/s_{\max}$ . Downstream of the bottleneck, every cluster still leaves as the same group of  $K$  messages, so the count-level branching structure is unchanged (the Borel cluster sizes, and the Fano factor at time scales beyond  $Ks_{\max}$ , are exactly those of the input), but the intra-cluster spacing has been regularised to  $s_{\max}$  and the instantaneous intensity is capped at  $1/s_{\max}$ . The clustering has not been removed; it has been made harmless to any consumer at least as fast as the bottleneck, which is why stages with service  $s_i \leq s_{\max}$  never queue. If the bottleneck is stage 1, all of the waiting in the pipeline sits in its mailbox and every downstream mailbox holds at most one message; if the bottleneck is stage  $b > 1$ , the stages before it queue partially (the raw burst is still bursty relative to their smaller services) and stage  $b$  absorbs the remainder, with the same total by Equation (4).

The practical consequences for the splitting recommendation are these. First, put the largest stage first, at ingress: end-to-end latency is unchanged by the order, but this confines all queueing to one mailbox, keeps the downstream mailboxes at depth one, and places the burst-absorbing queue where its receiver is continuously busy during a cluster (the regime in which the load-dependent hop cost of Section 4.7 is smallest). Second, any burst-aware logic, such as load shedding, coalescing of superseded book updates, or a cancel-on-burst risk rule, can only observe the raw burst upstream of the bottleneck; downstream it sees a rate- $1/s_{\max}$  stream and cannot tell a cluster from steady flow except by counting. Third, a packet-Hawkes fit run on the bottleneck's departure stream would report a lower intensity peak and a longer effective kernel while recovering the same branching ratio in counts; operators instrumenting a pipeline should fit on the ingress stream. Fourth, deterministic hop delay is what keeps downstream stages queue-free; hop jitter of amplitude  $\epsilon$  lets a downstream stage queue by at most  $\epsilon$  per message, so the reduction holds up to that jitter in practice.

**Fact** (Hawkes–Oakes 1974 cluster representation [2]). A stationary Hawkes process with background rate  $\mu > 0$  and branching ratio  $n \in [0, 1)$  is stochastically equivalent to a Poisson-cluster process: a rate- $\mu$  Poisson process of immigrant arrivals, each of which independently spawns an offspring cluster whose total size  $K$  (including the immigrant) follows the Borel distribution

$$\Pr(K = k) = \frac{(nk)^{k-1} e^{-nk}}{k!}, \quad k = 1, 2, 3, \dots \quad (7)$$

with mean  $E[K] = 1/(1 - n)$  and variance  $\text{Var}[K] = n/(1 - n)^3 [3]$ .

**Remark** (Where Hawkes enters: the size of the single-stage wait). Under Poisson arrivals at the utilisation of a CME packet feed ( $\rho = \bar{\lambda}_{\text{pkt}}T \ll 1$ ), the single-stage wait is zero for almost every message, so  $W_q^{(1)} = 0$  for every quantile  $q$  with  $\rho < 1 - q$  and Equation (6) has nothing to compress; that is Corollary 2, and at this corpus's rate it holds at  $p_{99}$  for all  $T \leq 16 \mu\text{s}$ . Under Hawkes arrivals the cluster representation puts  $K$  near-coincident messages behind one immigrant, and in the instantaneous-batch picture the last of them waits  $(K - 1)T$  on a single stage. Stirling's formula

applied to Equation (7) gives, for large  $k$ ,

$$\Pr(K = k) \approx \frac{1}{n\sqrt{2\pi}} k^{-3/2} (ne^{1-n})^k, \quad (8)$$

and  $ne^{1-n} < 1$  for every  $n < 1$ , with  $-\log(ne^{1-n}) \approx (1-n)^2/2$  near criticality. So the Borel batch size is light-tailed at every fixed  $n < 1$ : its survival function follows the critical  $k^{-1/2}$  power law only in the window  $k \ll k^* = (1-n)^{-2}$  and decays geometrically beyond it (Appendix A). The consequences for the paper are, first, that the largest clusters, and hence the single-stage cluster wait they produce, have a characteristic scale of order  $Tk^* = T(1-n)^{-2}$ , which at the corpus-median  $n \approx 0.8$  (Section 3.2) is  $\approx 25T$  and at  $n = 0.9$  would be  $100T$ , one to two orders of magnitude above the service floor and the hop cost; second, that heavy-tail results built on subexponential batch sizes, such as the single-big-batch principle for  $M^{[K]}/G/1$  [4] and the monotone-separable tail asymptotics of Baccelli et al. [14], do not apply to this batch distribution at any fixed  $n < 1$ , so no closed-form power-law exponent for the tail probability should be expected; and third, that none of this is needed for the design claim, because Theorem 1 applies the factor  $1/N$  to whatever  $W_q^{(1)}$  the feed produces. The size of  $W_q^{(1)}$  on real feeds is a measurement, and Part II reports it.

**Corollary 2** (Poisson null). *Under homogeneous Poisson arrivals at rate  $\bar{\lambda}$  with  $\rho = \bar{\lambda}T < 1 - q$ , the  $q$ -quantile of the single-stage wait is  $W_q^{(1)} = 0$ , because  $\Pr(W^{(1)} > 0) = \rho$  for  $M/D/1$  by PASTA. By Theorem 1,  $W_q^{(N)} \leq (N-1)h$ , and by Equation (4)  $W_q^{(N)} \geq (N-1)h$ , so*

$$W_q^{(N)} = (N-1)h \quad \text{for all } N \text{ and all } q \leq 1 - \rho. \quad (9)$$

At the utilisations of a CME packet feed (corpus-median 452 packets per second, so  $\rho = 0.0036$  at  $T = 8\mu s$ ) this covers  $p_{99}$  at every  $T \leq 16\mu s$  in the paper's sweep: end-to-end latency is  $T + (N-1)h$  at  $p_{50}$  and at  $p_{99}$  alike, and splitting costs  $(N-1)h$  with nothing to compress. Where  $\rho/N$  does exceed  $1 - q$ , Equation (6) still holds and the tandem compresses the small Poisson wait too — splitting then helps under Poisson as well, just for the ordinary utilisation reason rather than for the clustering reason this paper is about. The sweep contains six such  $p_{99}$  cells, all at the long-service end:  $(T, N) = (32, 1), (64, 1), (64, 2), (128, 1), (128, 2)$  and  $(128, 4)$ , spanning  $\rho/N = 0.0144$  to  $0.0578$ . It contains one such  $p_{95}$  cell ( $T = 128, N = 1, \rho = 0.0578 > 0.05$ ) and many  $p_{99.9}$  cells, and Section 4.1 shows the wait appearing in precisely those and in no others.

### 2.3 What the theory says, in plain terms

This subsection restates Theorems 1 and 1 and Corollary 2 without the formalism, for readers who build pipelines rather than analyse queues. The formal statements above remain the reference.

**Start from where a message's time goes.** A message's end-to-end latency is the sum of three things: the work done on it ( $T$  in total, however many stages that work is spread over), the hop tax ( $h$  for every mailbox boundary it crosses, so  $(N-1)h$  in an  $N$ -stage pipeline), and the time it spends sitting in queues waiting for a busy actor. The paper calls that third piece the wait,  $W$ . The work is fixed by the code and the hop tax is fixed by the framework; the wait is the only piece the architecture can change, and on a bursty feed it is the piece that dominates the tail.

**Proposition 1: one burst.** Suppose  $k$  messages land at once. On a single thread the last of them waits for the  $k-1$  ahead of it, each costing  $T$ : it waits  $(k-1)T$ . On an  $N$ -stage pipeline the first stage takes only  $T/N$  per message, so the last message waits  $(k-1)T/N$  at stage 1 and then flows

through the remaining stages without waiting again, because each of them is fed at exactly the rate it can serve. The pipeline clears a burst  $N$  times faster for the same total work, because  $N$  cores are working on  $N$  different messages at the same instant. That is what pipelining means, and it is the whole mechanism.

**Theorem 1: every message, whatever the traffic.** Two facts, each of which a programmer can check by hand.

1. In a pipeline where each stage does a fixed amount of work per message, only the slowest stage ever builds a queue. Every other stage receives messages no faster than it can process them, because the slowest stage upstream (or, if the slowest stage is downstream, the stages before it) meters them out. So the whole pipeline behaves exactly like one thread whose per-message work is the slowest stage’s work,  $s_{\max}$ , plus the hop tax. With equal stages,  $s_{\max} = T/N$ .
2. Take any recording of arrival times, real or synthetic, and replay it through two single-threaded servers: one doing  $T$  of work per message and one doing  $T/N$ . The second server’s queue wait for message number  $j$  is never more than  $1/N$  of the first server’s wait for the same message. The reason is that in the faster server every backlog is  $N$  times smaller and every idle gap between arrivals clears at least as much of it.

Put together: for every message in every replay, the pipeline’s wait is at most the single-thread wait divided by  $N$ , plus  $(N - 1)h$ . Because that holds message by message, it holds for whichever summary statistic the operator cares about: median,  $p_{99}$ ,  $p_{99.9}$ , or the worst case. And the bound is usually beaten. It is met exactly on an isolated burst that arrives all at once, and beaten whenever arrivals are spread out, because the pipeline’s first stage is idle more often than the single thread and uses each idle gap to drain more of its backlog. The paper writes the measured compression as  $N^{-\gamma}$ ; the theorem guarantees  $\gamma \geq 1$ , and Section 4 finds the bound beaten by a wide margin —  $\gamma = 3.3$  and  $2.7$  at  $T = 64$  and  $128 \mu\text{s}$ , and at shorter service times the tandem drives the excess to exactly zero, where no exponent is defined at all.

**Where the traffic model comes in.** Nowhere in the two theorems. The arrival pattern enters only through the size of the single-thread wait that the factor  $1/N$  is applied to. Under Poisson traffic at the load of a CME packet feed, the single thread is idle when almost every message arrives, so for tasks up to  $16 \mu\text{s}$  its  $p_{99}$  wait is zero; there is nothing to divide by  $N$ , and the pipeline is simply  $(N - 1)h$  slower (Corollary 2). Under the real feed, packets arrive in self-exciting clusters, the single thread’s  $p_{99}$  is 2.2 to 41.5 times  $T$  in the paper’s data for tasks of 16 to  $128 \mu\text{s}$ , and the pipeline divides the excess by  $N$  or better. The Hawkes cluster representation says why the clusters are large: a cluster’s size follows the Borel distribution, whose typical large cluster has of order  $(1 - n)^{-2}$  packets, about 25 at the corpus’s  $n \approx 0.8$ .

**The rules a designer can take away.**

- The tail improvement is set by the slowest stage,  $s_{\max}/T$ , and by nothing else about the decomposition. Balanced stages are the best case for a given number of hops.
- A cut that does not shrink the slowest stage costs one  $h$  at every quantile and improves nothing. Sub-tasks can be merged onto one thread until their combined work reaches  $s_{\max}$  at no tail cost.

- The order of the stages does not change any message’s latency. It changes only where the queue physically sits; putting the slowest stage first keeps all queueing in one mailbox.
- Every hop costs exactly  $h$  on every message, median and tail alike; there is no regime where the hop tax is hidden.
- Whether splitting is worth it comes down to one comparison: the single-thread wait at the quantile you care about, times  $(1 - s_{\max}/T)$ , against  $h$  per hop added. On a Poisson-like feed at HFT service times the left side is zero; on a clustered feed at the same rate it is tens to thousands of microseconds.

The guarantees above are exact inside a specific model, and a designer should know what that model leaves out. The rest of this section lists its assumptions and the questions it leaves to measurement.

**Fixed work per message.** Each stage is assumed to spend the same amount of time on every message. That is what makes only the slowest stage queue: a stage that sometimes takes longer than its upstream neighbour would occasionally fall behind and build its own queue, and then the pipeline’s wait would exceed the single-server wait at  $s_{\max}$ . The paper’s simulator uses constant service, so the theorems describe the simulator exactly; real actors have data-dependent work (a book update touching a deep level costs more than one touching the top), and for them the theorems are a model to be checked by measurement, which is the paper’s operational recommendation in any case (Section 8).

**No drops, no bounds, in-order.** Mailboxes are unbounded, messages are never discarded or coalesced, and each stage processes in arrival order. A framework that drops or merges messages under load changes the arithmetic (in the pipeline’s favour, at the cost of lost messages); a priority queue changes which message pays the wait but not the total.

**Fixed hop delay.** The hop is modelled as a constant  $h$ . A jittery hop lets a downstream stage queue by at most the jitter amplitude per message, so the reduction holds up to that amplitude. The paper’s calibration uses the low-load, wake-from-sleep value of  $h$ , which Section 4.7 argues is an upper bound inside a burst.

**Proposition 1 needs a dense burst.** The one-burst identity assumes the  $k$  messages arrive faster than the first stage can drain them, i.e. faster than one per  $T/N$ . For a burst that trickles in, the ratio in Equation (2) is  $kN/(k + N - 1 + N(N - 1)h/T)$ , which climbs toward  $N$  as the burst grows. Theorem 1 does not need this assumption; it covers every arrival pattern.


**What the theory does not predict.** It does not say how large the single-thread wait is on a given feed. That number depends on the feed’s clustering and utilisation and, as the remark after the Borel fact explains, has no clean closed form at a fixed branching ratio: the Borel cluster size is light-tailed for every  $n < 1$ , so the heavy-tail formulas that would give a power-law exponent do not apply. The paper measures the single-thread wait on real CME arrivals instead (Part II). Nor does the theory predict how far below the  $1/N$  bound the pipeline lands; that slack is also measured. What the theory does guarantee is that the pipeline can never do worse than  $1/N$  of the single-thread wait plus the hop tax, on any feed whatsoever, which is the statement a designer needs in order to trust the trade before measuring it.

## Part II

# Empirical experiment

### 3 Experiment setup

This section describes the empirical experiment that tests the analytical predictions of Section 2: the data corpus, the confirmation that packet arrivals are Hawkes-clustered (the reason the single-stage wait that Theorem 1 divides by  $N$  is large at all), the tandem-Lindley simulator, and the sweep design that produces the numbers reported in Section 4. Figure 1 shows the end-to-end data flow of a single 30-minute window; Algorithm 1 states the full per-window procedure in pseudocode.



**Figure 1.** Per-window data flow of the empirical experiment. A single session tape is chopped into 30-minute RTH windows. Each window’s messages are grouped by `packet_seq` into a real packet-arrival stream (Hawkes arm). A rate-matched Poisson-null arm is generated by uniformly shuffling the same arrival count over the same window. Both arms are run through the  $N$ -stage tandem Lindley simulator at  $N \in \{1, 2, 4, 8\}$ . Per-window quantiles are recorded and aggregated by corpus median across all 3512 windows in the 276-session corpus.

---

**Algorithm 1** Per-window procedure of the empirical experiment.

---

```
1: input: session tape  $S$ , service  $T$ , hop  $h$ , stage counts  $\mathcal{N}$ 
2: for each 30-min RTH window  $W$  in  $S$  with  $[T_0, T_1]$  do
3:    $\{m_i\} \leftarrow$  messages in  $W$ 
4:    $\{(\text{seq}, t)\} \leftarrow \{(\text{packet\_seq}(m_i), \text{sendingTime}(m_i))\}$ 
5:    $\text{arr}_H \leftarrow \text{sort}(\{\min_i t_i : \text{seq}_i = s\}_s)$   $\triangleright$  Hawkes arm (real arrivals)
6:    $n_{\text{pkt}}^\# \leftarrow |\text{arr}_H|$ 
7:    $\text{arr}_P \leftarrow \text{sort}(\text{Uniform}(\min \text{arr}_H, \max \text{arr}_H)^{n_{\text{pkt}}^\#})$   $\triangleright$  Poisson-null arm (seeded per window)
8:   for regime  $\in \{H, P\}$ ,  $N \in \mathcal{N}$  do
9:      $\text{lat} \leftarrow \text{TANDEMLINDLEY}(\text{arr}_{\text{regime}}, N, T, h)$ 
10:    record  $(S, W, \text{regime}, N, \{p_{50}, p_{95}, p_{99}, p_{99.9}, \max\}(\text{lat}))$ 
11:   end for
12: end for
13: return corpus-median of each recorded quantile across all  $(S, W)$  pairs
14: function TANDEMLINDLEY( $\text{arr}, N, T, h$ )
15:    $\text{dep} \leftarrow \text{arr}$ 
16:   for stage  $k = 1, \dots, N$  do
17:     if  $k > 1$  then  $\text{dep} \leftarrow \text{dep} + h$ 
18:   end if
19:    $\text{dep} \leftarrow \text{dep} + \text{LINDLEY}(\text{dep}, T/N)$ 
20: end for
21: return  $\text{dep} - \text{arr}$ 
22: end function
```

---

### 3.1 Data corpus

The corpus is 281 sessions of CME MDP3 channel 318 (NQ) front-month pcap-derived tapes covering 2025-01 through 2026-02, RTH only, holidays and half-days excluded. Each MBO record carries a `packet_seq` (the UDP packet sequence number) and two nanosecond timestamps: `transactTime`, the CME matching-engine event time carried in the MDP3 message body, and `sendingTime`, the gateway’s per-packet UDP header timestamp. Messages are grouped into packets by `packet_seq`, and a packet’s arrival time is its `sendingTime`.

The choice of clock matters, and `transactTime` is the wrong one. A single matching-engine transaction is routinely split across several UDP packets that leave the gateway microseconds apart, so all of those packets share one `transactTime`; keying arrivals on it collapses distinct packets onto a single instant and manufactures coincident arrivals that never occurred. Measured across 178.7 million packets, 3.18% carry a `transactTime` shared with at least one other packet, in groups of up to 66; the true emission spread within such a group has a median of  $15.4\mu\text{s}$  and reaches 51 ms. Because the mechanism under study is precisely how tightly packed arrivals queue behind one another, that artefact inflates every clustering and tail statistic downstream. `sendingTime` is emitted once per packet at nanosecond resolution and is the receiving pipeline’s actual input clock; over the same 178.7 million packets, ordered by `packet_seq`, it produces no ties and no sequence inversions whatsoever. All results below are keyed on `sendingTime`.

Each session tape is chopped into non-overlapping 30-minute RTH windows; for each window we compute the packet-arrival stream and the per-window quantities listed in Section 1.6. The corpus yields 3512 per-window records from the 276 sessions that contain at least one complete RTH window, roughly 13 windows per session; the remaining five of the 281 tapes are holiday or half-day

sessions with none.

### 3.2 Packet arrivals are Hawkes-clustered (empirical prerequisite)

The paper’s simulator input is the packet-arrival stream, so the first empirical question is whether that stream is clustered or Poisson. Three model-free and one model-based diagnostic on the 276-session, 3512-window NQ front-month corpus all reject Poisson decisively. The fraction of interarrival gaps shorter than one-tenth of the mean gap has a corpus median of 73% per 30-minute window, a 5th–95th percentile range of 54–78% and a full range of 38–85%, versus the Poisson benchmark  $1 - e^{-0.1} \approx 9.52\%$  that would obtain for any exponential-gap process regardless of rate; not one of the 3512 windows comes within a factor of three of the Poisson value. The Fano factor  $F(\tau) = \text{Var}[C_\tau]/\text{E}[C_\tau]$  of packet counts in bins of width  $\tau$  (written  $\tau$ , not  $T$ , to keep it clear of the service time) rises from  $F(0.1\text{ s}) = 51.7$  to  $F(60\text{ s}) = 1027$  across the corpus median, against  $F(\tau) \equiv 1$  for Poisson at every  $\tau$ ; the log-log fit gives a Hurst exponent  $H = 0.74$ , against  $H = 0.5$  for Poisson. These are within-window figures, corpus-median over the 3512 windows; pooling a whole session instead raises the Fano level by roughly a factor of four without changing the shuffle ratio, for reasons taken up in Section 6.4. Running the identical estimator over the paper’s own Poisson-null arm returns  $F(\tau) = 1.00$  at  $\tau = 0.1$  through 5 s, 0.94 at 60 s, and  $H = 0.50$ , which confirms that the estimator is unbiased at every bin width used here and that the Hawkes figures above are not an artefact of binning or of finite window length. Fitting the standard exponential-kernel Hawkes process  $\lambda(t) = \mu + \sum_{t_i < t} \alpha e^{-\beta(t-t_i)}$  to each window by MLE (Ogata log-likelihood with the standard recursive intensity, closed-form compensator, Nelder–Mead over  $(\mu, \alpha, \beta)$ ) gives a branching ratio  $n = \alpha/\beta$  with a corpus median of 0.798 and a 5th–95th percentile range of 0.761–0.832 across the 3512 windows (extremes 0.566 and 1.000, the latter a single window pinned at the fit boundary), with every fit converged; the same optimum is recovered by a reparameterised L-BFGS-B fit to four decimals on spot-checked windows. This is the same near-critical regime that Filimonov and Sornette [24] and Hardiman et al. [25] document on ES, at a somewhat lower level than their  $n \approx 0.9$ ; they fit trade arrivals rather than packet arrivals, and a packet aggregates several book events, which is expected to lower the fitted self-excitation. A Fisher–Yates shuffle of the interarrival gap sequence within each window, which preserves the marginal gap distribution exactly and destroys only the ordering, removes most but not all of the excess: the corpus-median Fano at  $\tau = 5\text{ s}$  falls from 259.7 to 17.6, and at  $\tau = 1\text{ s}$  from 108.4 to 17.5, i.e. 94% and 85% of the excess over the Poisson value of 1. What the shuffle removes is specifically the *scaling*: the shuffled Fano is flat at 16.1–17.6 across every bin width from 0.1 to 60 s, giving  $H = 0.50$  against the unshuffled  $H = 0.74$ . All of the long-range dependence therefore lives in the ordering, which is what distinguishes a self-exciting process from a renewal process. The flat residual of  $\approx 17$  is a scale-free over-dispersion contributed by the gap marginal itself, which is heavy-tailed; so the corpus is neither a renewal process nor a pure Hawkes process with light-tailed gaps, but a self-exciting process with heavy-tailed gaps. We report the split rather than the headline because it bounds what the Hawkes fit can be asked to carry: the fitted  $n$  above describes the ordering component, and the paper’s simulator is in any case driven by the recorded arrivals themselves, not by the fitted process. The paper’s downstream simulator is therefore driven by an arrival process that is empirically far from Poisson, and any tail-latency finding based on a matched-rate Poisson null represents a hard counterfactual, not a straw man. We use the Hawkes fit only as a parsimonious description of the observed clustering; the tandem reduction of Section 2 is arrival-law agnostic and does not depend on the Hawkes assumption — any arrival process with the same clustering drives the same single-server queue that the bound compresses, which is precisely why the model-free diagnostics above, not the Hawkes fit, carry the empirical premise.

### 3.3 Simulator

Packet arrivals  $t_1 < t_2 < \dots$  enter stage 1. Each stage  $k$  is a single-server FIFO queue with deterministic service  $T/N$  (own actor, own thread), and departures from stage  $k$  are delayed by  $h$  before arriving at stage  $k + 1$ . Standard Lindley recursion per stage:

$$W_{i+1}^{(k)} = \max\left(0, W_i^{(k)} + \frac{T}{N} - (a_{i+1}^{(k)} - a_i^{(k)})\right), \quad (10)$$

with  $a_i^{(k+1)} = a_i^{(k)} + W_i^{(k)} + T/N + h$ .

### 3.4 Pipelining is what makes multi-stage work

The tandem is a *pipelined* design: each stage’s actor runs on its own thread, so at the same wall-clock instant stage 2 can be processing message  $k - 1$  while stage 1 is processing message  $k$ . The simulator captures this implicitly because Lindley is applied separately per stage and each stage tracks its own “when did the previous message leave” timeline — the two servers’ busy intervals overlap on the wall-clock axis.

Worked example. Three packets arriving at  $t = 0, 1, 2 \mu\text{s}$ ; single-stage total service  $T = 8 \mu\text{s}$  (a service floor in the sweep of Table 1); two-stage tandem with per-stage service  $T/2 = 4 \mu\text{s}$  and hop  $h = 1.7 \mu\text{s}$ :

|                                                  | msg 1      | msg 2       | msg 3       |
|--------------------------------------------------|------------|-------------|-------------|
| single-stage arrival                             | 0          | 1           | 2           |
| single-stage latency ( $\mu\text{s}$ )           | 8.0        | 15.0        | <b>22.0</b> |
| tandem arrive stage 1                            | 0.0        | 1.0         | 2.0         |
| tandem wait stage 1                              | 0.0        | 3.0         | 6.0         |
| tandem leave stage 1                             | 4.0        | 8.0         | 12.0        |
| tandem arrive stage 2 (+ $h$ )                   | 5.7        | 9.7         | 13.7        |
| tandem wait stage 2                              | 0.0        | 0.0         | 0.0         |
| tandem leave stage 2                             | 9.7        | 13.7        | 17.7        |
| <b>tandem latency (<math>\mu\text{s}</math>)</b> | <b>9.7</b> | <b>12.7</b> | <b>15.7</b> |

At wall-clock  $t \approx 7$ , stage 1’s actor is busy on message 2 while stage 2’s actor is busy on message 1 — two actors doing useful work at the same time. Message 1 pays a  $1.7 \mu\text{s}$  hop cost (latency  $8.0 \rightarrow 9.7$ ); message 3, at the burst tail, drops from 22.0 to  $15.7 \mu\text{s}$ . Median hurts, tail wins — same message count, throughput  $2/T$  rather than  $1/T$ . This is exactly the mechanism of Proposition 1.

**Pipelining is a necessary condition, not a nice-to-have.** The tail-latency improvement in Proposition 1 and Theorem 1 comes entirely from concurrent execution across the  $N$  stages: at any wall-clock instant, several different messages are simultaneously in flight at different stages of the pipeline, and it is that overlap that raises system throughput from  $1/T$  to  $\min_i 1/s_i$  (equal to  $N/T$  under equal-sized stages). Take away the concurrency — for instance by chaining  $N$  synchronous handoffs on a single thread, where the sender blocks until each downstream stage finishes — and the entire benefit disappears. In that case each message serially traverses all  $N$  stages before the next message can start; total service work per message is still  $T$ , so throughput remains  $1/T$ , and the pipeline adds  $(N - 1)h$  of hop cost with no compensating tail win. Under Hawkes bursts a synchronous  $N$ -stage chain is strictly worse than a single-stage event loop by exactly  $(N - 1)h$  per

message. Pipelining — concurrent execution of different messages at different stages on different threads — is therefore not an optimisation on top of the tandem design, it is the design; the paper’s simulator captures this by applying Lindley separately per stage, giving each stage its own busy-idle timeline. The framework-implementation Section 8 discusses which KASPAR-HFT dispatch paths actually deliver this concurrency (the asynchronous BQueue path) and which do not (the synchronous `fast_send` path, which despite its low per-call cost does not support pipelining).

**Calibration and cost.** The paper reports numerical results at  $h = 1.7 \mu\text{s}$  — the one-way async mailbox cost obtained by halving the round-trip figure in Table 4 of Maciejewski [32] — swept across  $T \in \{2, 4, 8, 16, 32, 64, 128\} \mu\text{s}$  to cover the sub-hop regime through the slow-servicing regime. The simulator is deterministic given the tape and  $(T, h, N)$ ; one  $O(n)$  pass per stage per window makes its wall-clock cost negligible next to the tape read.

### 3.5 Sweep design

The design asks whether an  $N$ -stage tandem reduces  $p_{99}$  latency under real Hawkes packet arrivals and whether the same architecture achieves the same compression under a rate-matched Poisson null. For each 30-minute window we run the tandem Lindley of Section 3.3 at  $N \in \{1, 2, 4, 8\}$  under two arrival regimes: the real packet-arrival timestamps from the tape (Hawkes arm) and a uniform-shuffle of the same  $n_{\text{pkt}}^{\#}$  arrivals over the same support (Poisson-null arm). The support is the observed packet span  $[\min \text{arr}_H, \max \text{arr}_H]$  rather than the nominal window  $[T_0, T_1]$ , so the null is rate-matched to the arrivals it replaces rather than to wall-clock; the two coincide on busy windows and the former is the stricter counterfactual on sparse ones. The null is therefore a *conditional homogeneous-Poisson* process: a homogeneous Poisson process at rate  $n_{\text{pkt}}^{\#}/(\max \text{arr}_H - \min \text{arr}_H)$  conditioned on the observed count over the observed support. Per (window,  $N$ , regime) we record  $p_{50}, p_{95}, p_{99}, p_{99.9}, \max$  of end-to-end latency. Aggregation is corpus-median across all windows; the two headline curves per scenario are  $p_{99}(N)$  and  $p_{50}(N)$  under Hawkes and Poisson overlaid.

The paper’s empirical calibration is  $h = 1.7 \mu\text{s}$  (one-way async cross-thread mailbox cost, per Maciejewski [32]’s Table 4 halved for one-way) swept over  $T \in \{2, 4, 8, 16, 32, 64, 128\} \mu\text{s}$  so the scenarios span from the sub-hop regime through the slow-servicing regime. Theorem 1 predicts  $p_{99}^H(N) - (N - 1)h - T \leq (p_{99}^H(1) - T)/N$  on every window, i.e. a fitted exponent  $\gamma \geq 1$  in  $\Delta(N) \approx \Delta(1)N^{-\gamma}$ ; Corollary 2 predicts  $p_{99}^P(N) = p_{50}^P(N) = T + (N - 1)h$  exactly in every cell whose per-stage utilisation satisfies  $\rho/N < 0.01$ , and a strictly positive Poisson wait in every cell where it does not; the prediction is two-sided, so a Poisson  $p_{99}$  sitting on the floor at  $\rho/N > 0.01$  would falsify it just as a Poisson  $p_{99}$  above the floor at  $\rho/N < 0.01$  would. Any of (i) a window on which the Hawkes bound is violated, (ii) a Poisson  $p_{99}(N)$  on the wrong side of the  $\rho/N = 0.01$  boundary, or (iii) Hawkes  $p_{50}(N)$  growing at slope  $\gg h$  per stage falsifies the paper’s model of the simulator.


## 4 Empirical results

### 4.1 Tail collapse under Hawkes, and its near-absence under Poisson at HFT service times

Table 1 reports corpus-median end-to-end latency quantiles for every  $(T, N)$  cell of the sweep at  $h = 1.7 \mu\text{s}$ , on 276 sessions and 3512 thirty-minute windows (five of the 281 session tapes contain no complete RTH window and are dropped). Figure 2 plots the same numbers. The Poisson-null columns are the uniform-shuffle counterfactual of Section 3.5 with the same packet count per window.

| $T$ | $N$ | Hawkes (real arrivals) |          |            | Poisson null |          |            |
|-----|-----|------------------------|----------|------------|--------------|----------|------------|
|     |     | $p_{50}$               | $p_{99}$ | $p_{99.9}$ | $p_{50}$     | $p_{99}$ | $p_{99.9}$ |
| 2   | 1   | 2.00                   | 2.00     | 2.00       | 2.00         | 2.00     | 2.00       |
| 2   | 2   | 3.70                   | 3.70     | 3.70       | 3.70         | 3.70     | 3.70       |
| 2   | 4   | 7.10                   | 7.10     | 7.10       | 7.10         | 7.10     | 7.10       |
| 2   | 8   | 13.90                  | 13.90    | 13.90      | 13.90        | 13.90    | 13.90      |
| 4   | 1   | 4.00                   | 4.00     | 4.00       | 4.00         | 4.00     | 5.78       |
| 4   | 2   | 5.70                   | 5.70     | 5.70       | 5.70         | 5.70     | 5.70       |
| 4   | 4   | 9.10                   | 9.10     | 9.10       | 9.10         | 9.10     | 9.10       |
| 4   | 8   | 15.90                  | 15.90    | 15.90      | 15.90        | 15.90    | 15.90      |
| 8   | 1   | 8.00                   | 8.61     | 10.14      | 8.00         | 8.00     | 13.79      |
| 8   | 2   | 9.70                   | 9.70     | 9.70       | 9.70         | 9.70     | 11.48      |
| 8   | 4   | 13.10                  | 13.10    | 13.10      | 13.10        | 13.10    | 13.10      |
| 8   | 8   | 19.90                  | 19.90    | 19.90      | 19.90        | 19.90    | 19.90      |
| 16  | 1   | 16.00                  | 34.59    | 53.97      | 16.00        | 16.00    | 29.84      |
| 16  | 2   | 17.70                  | 18.31    | 19.84      | 17.70        | 17.70    | 23.49      |
| 16  | 4   | 21.10                  | 21.10    | 21.10      | 21.10        | 21.10    | 22.88      |
| 16  | 8   | 27.90                  | 27.90    | 27.90      | 27.90        | 27.90    | 27.90      |
| 32  | 1   | 32.00                  | 146.14   | 259.11     | 32.00        | 42.02    | 62.02      |
| 32  | 2   | 33.70                  | 52.29    | 71.67      | 33.70        | 33.70    | 47.54      |
| 32  | 4   | 37.10                  | 37.71    | 39.24      | 37.10        | 37.10    | 42.89      |
| 32  | 8   | 43.90                  | 43.90    | 43.90      | 43.90        | 43.90    | 45.68      |
| 64  | 1   | 98.88                  | 820.70   | 1548.68    | 64.00        | 106.73   | 126.72     |
| 64  | 2   | 65.70                  | 179.84   | 292.81     | 65.70        | 75.72    | 95.72      |
| 64  | 4   | 69.10                  | 87.69    | 107.07     | 69.10        | 69.10    | 82.94      |
| 64  | 8   | 75.90                  | 76.51    | 78.04      | 75.90        | 75.90    | 81.69      |
| 128 | 1   | 364.45                 | 5316.85  | 12487.72   | 128.00       | 237.65   | 289.07     |
| 128 | 2   | 164.58                 | 886.40   | 1614.38    | 129.70       | 172.43   | 192.42     |
| 128 | 4   | 133.10                 | 247.24   | 360.21     | 133.10       | 143.12   | 163.12     |
| 128 | 8   | 139.90                 | 158.49   | 177.87     | 139.90       | 139.90   | 153.74     |

**Table 1.** Corpus-median end-to-end latency quantiles in  $\mu\text{s}$  at  $h = 1.7 \mu\text{s}$ , for each service floor  $T$  ( $\mu\text{s}$ ) and stage count  $N$ . 276 sessions, 3512 thirty-minute RTH windows of CME MDP3 NQ front-month packet arrivals. The Poisson null is the uniform shuffle of the same arrival count per window.



**Figure 2.** Corpus-median latency versus stage count  $N$  for each service floor  $T$  (one shade per  $T$ , each curve labelled at its  $N = 1$  end). Left: median. Centre:  $p_{99}$ , Hawkes solid and Poisson null dashed. Right: the Hawkes tail excess ratio  $\Delta(N)/\Delta(1)$  against the Theorem 1 bound  $1/N$  (dotted); open markers are cells where the tandem drives  $\Delta(N)$  to exactly zero, drawn at the panel floor because zero cannot be placed on a log axis. The median rises by exactly  $h$  per hop wherever the queue is empty at the median, which is every cell except the three longest-service Hawkes cells. The Hawkes  $p_{99}$  falls with  $N$  only at  $T \geq 16 \mu s$ ; at  $T \leq 8 \mu s$  there is no tail to compress and it rises with the hop count, tracking the median. The Poisson  $p_{99}$  likewise tracks the median up to  $T = 16 \mu s$  and separates from it at  $T \geq 32 \mu s$ , where utilisation queueing begins.

Under the Poisson null the queue is empty when almost every message arrives, so every reported quantile equals the service floor,  $p_{50} = p_{99} = T + (N - 1)h$ , in every cell with  $T \leq 16 \mu s$ . The first exception is  $T = 32 \mu s$ ,  $N = 1$ , where Poisson  $p_{99}$  is  $42.0 \mu s$  against a floor of 32: at the corpus-median packet rate of 452 packets per second the utilisation there is  $\rho = 452 \times 32 \mu s = 0.0145$ , above the 0.01 at which Corollary 2 says the Poisson wait must first appear at  $p_{99}$ , and it does; at  $N = 2$  the per-stage utilisation halves to 0.007 and the wait vanishes again. Beyond that the null arm is a perfectly ordinary queue: at  $T = 64$  and  $128 \mu s$  the single-stage Poisson  $p_{99}$  reaches 106.7 and  $237.7 \mu s$ , and splitting reduces it, to 75.7 and  $172.4 \mu s$  at  $N = 2$ . *Under Poisson, splitting is not useless in general — it is useless in the regime an HFT hot path occupies.* The  $p_{99.9}$  column tells the same story one decade down: the Poisson wait is visible at  $N = 1$  for every  $T \geq 4 \mu s$  (single-stage  $\rho > 10^{-3}$ ) and at  $N = 2$  for  $T \geq 8 \mu s$ , and absent wherever the per-stage  $\rho/N$  falls below  $10^{-3}$ . Every one of these switch-on points is the  $\rho = 1 - q$  boundary of Corollary 2, which the corpus reproduces to three significant figures at  $q = 0.95$ ,  $0.99$  and  $0.999$  (Section 4.6). The null arm is therefore not inert; it is a queue with a small, exactly predictable wait, and over  $T \leq 16 \mu s$  splitting under Poisson loses  $(N - 1)h$  at every quantile the operator would look at.

Under Hawkes the story is different, but — and this is the central quantitative correction the packet-clock rerun makes to an earlier draft of this work — only above a sharp threshold. At  $T = 2$  and  $4 \mu s$  there is no tail at all: the single-stage  $p_{99}$  and  $p_{99.9}$  both sit exactly on the service floor, the clustered feed is indistinguishable from the Poisson null, and splitting buys nothing while costing  $(N - 1)h$ . The tail appears between  $T = 4$  and  $8 \mu s$  and then grows explosively: single-stage  $p_{99}$  is 1.00, 1.00, 1.08, 2.16, 4.57, 12.8 and 41.5 times the service floor at  $T = 2, 4, 8, 16, 32, 64, 128 \mu s$ , and single-stage  $p_{99.9}$  reaches  $97.6 T$  at  $T = 128$ . Splitting into two stages cuts the  $p_{99}$  from 34.59 to  $18.31 \mu s$  at  $T = 16$ , from 146.1 to  $52.3 \mu s$  at  $T = 32$ , from 820.7 to  $179.8 \mu s$  at  $T = 64$ , and from 5316.9 to  $886.4 \mu s$  at  $T = 128$ , in each case for a median cost of  $1.7 \mu s$  per hop; the two-stage  $p_{99}$  is below the single-stage  $p_{99}$  on 100% of the 3512 windows at every  $T \geq 16 \mu s$ . The stage count that minimises  $p_{99}$  rises with  $T$ :  $N = 1$  at  $T \leq 8$ ,  $N = 2$  at  $T = 16$ ,  $N = 4$  at  $T = 32$ ,  $N = 8$  at  $T = 64$  and  $128$ . The  $T = 2$  and  $4 \mu s$  rows, both below the floor, are in the sweep to demonstrate that boundary, and they do: splitting there is strictly dominated under both arrival laws. Section 4.5 resolves the  $4-8 \mu s$

onset to  $1\text{ }\mu\text{s}$  and Section 4.6 explains what sets its location.

## 4.2 Median cost is linear in $N$ , exactly as $(N - 1)h$

The  $p_{50}$  columns of Table 1 match the single-message formula of Proposition 1

$$p_{50}(N) = T + (N - 1)h \quad \text{with } h = 1.7\text{ }\mu\text{s},$$

in all 28 cells of the Poisson arm and in 25 of the 28 Hawkes cells — for instance 8.00, 9.70, 13.10, 19.90 at  $T = 8$ . Under Poisson the identity holds because the queue is empty; under Hawkes it holds because the median message arrives in a quiet inter-cluster stretch where the queue is also empty. The hop tax is therefore exactly  $h$  per boundary on the median, with no hidden component, as Theorem 1 says it must be at every quantile.

The three exceptions are informative and are all at the far end of the sweep:  $(T, N) = (64, 1)$ ,  $(128, 1)$  and  $(128, 2)$ , where the Hawkes  $p_{50}$  is 98.88, 364.45 and  $164.58\text{ }\mu\text{s}$  against floors of 64.00, 128.00 and 129.70. There the server is slow enough that the median message no longer lands in a quiet stretch — the clusters have merged into a persistently backlogged queue and even the typical message waits. This is the regime boundary showing up in the median rather than in the tail, and it is a boundary the paper’s recommendations stay well clear of: an HFT stage with a  $128\text{ }\mu\text{s}$  service time on this feed is not a stage to pipeline, it is a stage to rewrite. Note that splitting repairs the identity as soon as the per-stage service drops back under the threshold — at  $T = 128$  the median is back on the floor by  $N = 4$ .

## 4.3 Design equation

From the corpus:

$$p_{99}^{\text{Hawkes}}(N) = T + (N - 1)h + \Delta(N), \quad (11)$$

$$p_{50}^{\text{Hawkes}}(N) = T + (N - 1)h, \quad (12)$$

where  $\Delta(N) \equiv p_{99}^{\text{Hawkes}}(N) - p_{50}^{\text{Hawkes}}(N)$  is the “clustering tail excess”. Table 2 lists it, together with the ratio  $\Delta(N)/\Delta(1)$  that Theorem 1 bounds by  $1/N$ .

| $T$   | $\Delta(N)$ ( $\mu\text{s}$ ) |        |        |       | $\Delta(N)/\Delta(1)$ |       |       | $\gamma$ |
|-------|-------------------------------|--------|--------|-------|-----------------------|-------|-------|----------|
|       | $N=1$                         | $N=2$  | $N=4$  | $N=8$ | $N=2$                 | $N=4$ | $N=8$ |          |
| 2     | 0.00                          | 0.00   | 0.00   | 0.00  | —                     | —     | —     | —        |
| 4     | 0.00                          | 0.00   | 0.00   | 0.00  | —                     | —     | —     | —        |
| 8     | 0.61                          | 0.00   | 0.00   | 0.00  | 0.000                 | 0.000 | 0.000 | —        |
| 16    | 18.59                         | 0.61   | 0.00   | 0.00  | 0.033                 | 0.000 | 0.000 | —        |
| 32    | 114.14                        | 18.59  | 0.61   | 0.00  | 0.163                 | 0.005 | 0.000 | —        |
| 64    | 721.82                        | 114.14 | 18.59  | 0.61  | 0.158                 | 0.026 | 0.001 | 3.32     |
| 128   | 4952.41                       | 721.82 | 114.14 | 18.59 | 0.146                 | 0.023 | 0.004 | 2.68     |
| bound |                               |        |        |       | 0.500                 | 0.250 | 0.125 | $\geq 1$ |

**Table 2.** Corpus-median clustering tail excess  $\Delta(N) = p_{99}(N) - p_{50}(N)$  under Hawkes arrivals at  $h = 1.7\text{ }\mu\text{s}$ , its ratio to the single-stage value, and the least-squares exponent  $\gamma$  in  $\Delta(N) \approx \Delta(1)N^{-\gamma}$ . The last row is the Theorem 1 bound.

Two features of Table 2 are the theorems made visible. First,  $\Delta$  depends on  $(T, N)$  only through the per-stage service  $T/N$ : every diagonal of constant  $T/N$  carries one number — 0.00 at  $T/N \leq 4 \mu\text{s}$ , 0.61 at  $T/N = 8$ , 18.59 at 16, 114.14 at 32, 721.82 at 64 — to the reported precision, across up to four different  $(T, N)$  pairs, with no fitting of any kind. That is the reduction identity Equation (4) on real arrivals: the tandem’s tail is the single-server tail at  $s_{\max}$ , and the table is really a table of  $\Delta_{\text{single}}(s)$  at  $s \in \{0.25, 0.5, \dots, 64, 128\} \mu\text{s}$  read off the diagonals. The reduction is the single most directly verifiable claim in the paper and the data confirm it exactly.

Second, every ratio sits at or below the bound  $1/N$ , and the pathwise check of Section 5 finds no message on any window above it — but the margin is now enormous rather than tight, and the reason is the threshold. Because  $\Delta_{\text{single}}(s)$  is identically zero for  $s \lesssim 7.5 \mu\text{s}$  (Section 4.5), a split that pushes the per-stage service below that floor does not merely divide the tail excess, it removes it. At  $T = 32$ ,  $\Delta$  falls  $114.14 \rightarrow 18.59 \rightarrow 0.61 \rightarrow 0.00$  as  $N$  goes  $1 \rightarrow 2 \rightarrow 4 \rightarrow 8$ : a ratio of 0.005 at  $N = 4$  against a bound of 0.25. The fitted exponent  $\gamma$  in  $\Delta(N) \approx \Delta(1)N^{-\gamma}$  is correspondingly steep, 3.32 at  $T = 64$  and 2.68 at  $T = 128$ , well above the  $\gamma = 1$  that the bound alone would permit; at  $T \leq 32$  the tandem drives  $\Delta(8)$  to exactly zero and no power law is fitted, since the least-squares slope would be reading the numerical floor rather than the data. The excess decay is therefore not a  $1/N$  division at all in this corpus — it is a division that terminates in a cliff, and the design consequence is that the operator’s goal is not to make  $N$  large but to make  $s_{\max}$  land under the feed’s tight-gap floor. The utilisation slack described after Theorem 1 contributes as well, and Appendix B reproduces it on synthetic Borel clusters.

The operator’s stage-count optimum is

$$N^* = \arg \min_N \left[ (N-1)h + \kappa \Delta(N) \right]$$

for the operator’s tail weight  $\kappa$  (the number of microseconds of median the operator will spend to remove one microsecond of  $p_{99}$  excess). Over the simulated grid  $N \in \{1, 2, 4, 8\}$  at  $h = 1.7 \mu\text{s}$  the breakpoints are:

| $T$ ( $\mu\text{s}$ ) | $N^* = 1$         | $N^* = 2$                     | $N^* = 4$                    | $N^* = 8$           |
|-----------------------|-------------------|-------------------------------|------------------------------|---------------------|
| 2                     | always            | —                             | —                            | —                   |
| 4                     | always            | —                             | —                            | —                   |
| 8                     | $\kappa < 2.8$    | $\kappa \geq 2.8$             | —                            | —                   |
| 16                    | $\kappa < 0.095$  | $0.095 \leq \kappa < 5.6$     | $\kappa \geq 5.6$            | —                   |
| 32                    | $\kappa < 0.018$  | $0.018 \leq \kappa < 0.19$    | $0.19 \leq \kappa < 11$      | $\kappa \geq 11$    |
| 64                    | $\kappa < 0.0028$ | $0.0028 \leq \kappa < 0.036$  | $0.036 \leq \kappa < 0.38$   | $\kappa \geq 0.38$  |
| 128                   | $\kappa < 0.0004$ | $0.0004 \leq \kappa < 0.0056$ | $0.0056 \leq \kappa < 0.071$ | $\kappa \geq 0.071$ |

A dash means that stage count is never optimal at any  $\kappa$ : at  $T \leq 8 \mu\text{s}$  the tail excess is already zero by  $N = 2$ , so further splitting can only add hop cost. An operator who values tail and median equally ( $\kappa = 1$ ) lands on  $N^* = 1$  at  $T = 2, 4$  and 8, 2 at  $T = 16$ , 4 at  $T = 32$ , and 8 at  $T = 64$  and  $128 \mu\text{s}$ , which is also the ordering of the raw  $p_{99}$  minima in Table 1. The practical reading is that over the service times an HFT hot path actually occupies, an operator has to weight the tail several times more heavily than the median before splitting pays at all — and once past the threshold, has to weight it barely at all.

At real cut points the stages are unequal, and by Remark 2.2 the tail excess depends on the partition  $P$  of the servicing chain only through its largest stage,  $\Delta(P) = \Delta_{\text{single}}(s_{\max}(P))$ , i.e. the single-stage tail excess measured at service  $s_{\max}$ . The optimum over partitions is then

$$P^* = \arg \min_P \left[ \text{hops}(P)h + \kappa \Delta_{\text{single}}(s_{\max}(P)) \right],$$

which Table 2 lets the operator evaluate directly, since its diagonals are  $\Delta_{\text{single}}(s)$  on a power-of-two grid of  $s$ . Two rules of thumb fall out. A cut is worth its hop only if it lowers  $s_{\text{max}}$  by enough that  $\kappa[\Delta_{\text{single}}(s_{\text{max}}) - \Delta_{\text{single}}(s'_{\text{max}})] > h$ ; and among partitions with the same  $s_{\text{max}}$  the one with the fewest hops wins outright, so sub-tasks downstream of the bottleneck should be merged onto one thread until their sum reaches  $s_{\text{max}}$ .

**Design principle (only  $s_{\text{max}}$  matters).** For a servicing chain cut at fixed boundaries, the tail-latency question is not “how many threads?” but “what is the smallest achievable largest stage  $s_{\text{max}}$ ?” By Theorem 1 the clustering tail excess of any partition is  $\Delta_{\text{single}}(s_{\text{max}})$ , independent of the stage count; a cut that does not lower  $s_{\text{max}}$  compresses no tail and still costs one hop  $h$ . Choose the partition with the smallest reachable  $s_{\text{max}}$ , then the fewest hops that realise it.

**Floor on  $T$ : short tasks cannot be usefully split.** Equations (11)–(12) carry an obvious lower bound on the tasks the paper’s recommendation applies to. Splitting a task of total service  $T$  into  $N$  stages spends  $(N - 1)h$  at every quantile to compress a tail excess  $\Delta(1)$  that is itself a strongly increasing function of  $T$  (0 at  $T \leq 4\mu\text{s}$ , rising to  $39T$  at  $T = 128$ , Table 2). When  $T$  is at or below the feed’s tight-gap floor there is no tail excess to remove at all, and the design equation prescribes  $N^* = 1$  mechanically; a hop there is pure cost. The sweep shows exactly this: at  $T = 2$  and  $4\mu\text{s}$ ,  $\Delta(1) = 0$ , so a split removes nothing and  $p_{99}$  rises by the full hop, from 2.00 to 3.70 and from 4.00 to 5.70. Above the floor the arithmetic reverses sharply: at  $T = 16\mu\text{s}$ ,  $\Delta(1) = 18.59\mu\text{s}$ , of which the split removes 17.98 for the same 1.7, and  $p_{99}$  falls from 34.59 to 18.31 on every window.

Two separate necessary conditions are at work here and it is worth keeping them apart. The first is a property of the *framework*: a split cannot pay unless the tail excess it removes exceeds the hop cost it adds, which on this framework’s shipping BQueue path ( $h \approx 1.7\mu\text{s}$ ) puts a floor at roughly  $T \gtrsim 2h \approx 3.4\mu\text{s}$ . The second is a property of the *feed*: there is no tail excess to remove at all until  $T$  exceeds the arrival stream’s tight-gap floor, about  $7.5\mu\text{s}$  on this corpus (Section 4.6). The binding constraint is whichever is larger, and on CME MDP3 channel 318 with a BQueue mailbox it is the feed’s, by roughly a factor of two — which is why  $T = 4\mu\text{s}$ , though comfortably above  $2h$ , still has  $\Delta(1) = 0$  and must not be split. A framework with a much more expensive mailbox would find its own floor binding instead. Throughout the rest of the paper, “the floor” means the larger of the two. Sub-microsecond tasks, for instance a single SBE integer decode or a small handler that fires from a strategy fast path, belong on `fast_send` rather than on a cross-thread mailbox, regardless of feed statistics.

#### 4.4 Across windows

The corpus medians understate the tail on the busiest windows. Across the 3512 windows the single-stage  $p_{99}/T$  has a median of 1.0, 1.0, 1.1, 2.2, 4.6, 12.8, 41.5 and a 95th percentile of 1.0, 1.0, 1.1, 2.5, 6.1, 19.9, 81.7 at  $T = 2, 4, 8, 16, 32, 64, 128\mu\text{s}$ , so one window in twenty at  $T = 32$  has a single-stage  $p_{99}$  above  $195\mu\text{s}$  and at  $T = 128$  above 10.5ms. The tandem’s win is not confined to the median window:  $p_{99}(2) < p_{99}(1)$  and  $p_{99}(4) < p_{99}(1)$  both hold on 100% of windows at every  $T \geq 16$ , and on no window at all at  $T \leq 8$ , where there is no tail to compress and the split pays only hop cost. The crossover is sharp and sits just below  $T = 10\mu\text{s}$ : on the refined sweep of Section 4.5,  $p_{99}(2) < p_{99}(1)$  on 0% of windows at  $T = 8\mu\text{s}$ , 67.0% at  $T = 9\mu\text{s}$  and 99.7% at  $T = 10\mu\text{s}$ . The per-window fitted branching ratio has median 0.798 and a 5th–95th percentile range of 0.761–0.832 (Section 3.2); how the tail excess depends on it, and on the window’s packet rate, is taken up in Section 4.6, where the answer turns out to be service-time-dependent and not the one

we expected. The intervening  $p_{95}$  quantile is recorded per window in the sweep (Section 3.5) and omitted from the tables above only for brevity; it lies between the reported  $p_{50}$  and  $p_{99}$  in every cell by construction, and under the Poisson null it equals the service floor wherever  $p_{99}$  does, since  $\rho/N < 1 - 0.99$  implies  $\rho/N < 1 - 0.95$ . The 3512 windows are treated as evaluation units for the descriptive statistics reported here; they are nested within 276 sessions and are not independent, so any statistical uncertainty on the aggregates would be clustered at the session level rather than treating 3512 as an independent sample size.

#### 4.5 Where the tail switches on

The coarse sweep brackets the onset of the clustering tail between  $T = 8$  and  $T = 16 \mu\text{s}$ , which is too wide to turn into a design rule. We therefore reran the single- and two-stage arms on the same 3512 windows at  $T \in \{5, 6, 7, 8, 9, 10, 12, 14, 16\} \mu\text{s}$ . The onset is far sharper than the coarse grid suggests:

| $T \ (\mu\text{s})$ | $p_{99}$ Hawkes | excess over $T$ | windows with a tail | $p_{99}(2) < p_{99}(1)$ |
|---------------------|-----------------|-----------------|---------------------|-------------------------|
| 5                   | 5.00            | 0.00            | 0.0%                | 0.0%                    |
| 6                   | 6.00            | 0.00            | 0.0%                | 0.0%                    |
| 7                   | 7.00            | 0.00            | 4.5%                | 0.0%                    |
| 8                   | 8.61            | 0.61            | 98.6%               | 0.0%                    |
| 9                   | 10.85           | 1.85            | 100.0%              | 67.0%                   |
| 10                  | 13.52           | 3.52            | 100.0%              | 99.7%                   |
| 12                  | 19.83           | 7.83            | 100.0%              | 100.0%                  |
| 14                  | 26.68           | 12.68           | 100.0%              | 100.0%                  |
| 16                  | 34.59           | 18.59           | 100.0%              | 100.0%                  |

Two distinct thresholds are visible, and conflating them is the easiest mistake to make with this data. The first is the *tail* threshold: between  $T = 7$  and  $T = 8 \mu\text{s}$  the fraction of windows with any  $p_{99}$  excess jumps from 4.5% to 98.6%. The second is the *splitting* threshold, about  $1.5 \mu\text{s}$  later, between  $T = 8$  and  $T = 10 \mu\text{s}$ : a tail must exist *and* exceed the hop cost before a split repays it. At  $T = 8 \mu\text{s}$  nearly every window has a tail, but its median excess is  $0.61 \mu\text{s}$ , well under  $h = 1.7 \mu\text{s}$ , and splitting loses on every window. By  $T = 10 \mu\text{s}$  the excess is  $3.52 \mu\text{s}$ , comfortably above  $h$ , and splitting wins on 99.7%.

The location of the first threshold is not arbitrary. It is set by the feed’s *tight-gap* structure rather than by its average rate — a distinction that the next section shows is the key to the whole regime.

#### 4.6 Conditioning on intensity and criticality

Everything reported so far is conditioned on the service time  $T$  alone. That is not obviously enough. A Hawkes tail is generated by two things the corpus median hides: how busy the window is, measured by the mean packet rate  $\bar{\lambda}$ , and how close to critical its self-excitation is, measured by the branching ratio  $n = \alpha/\beta$ . The natural expectation — ours, before measuring — is that the tail, and therefore the case for splitting, strengthens with both, so that a rule stated in  $T$  alone would be a rule that silently averages over market states and would mislead on the busy ones. We report the check because the expectation turns out to be wrong over the range of service times an HFT hot path actually occupies, and a rule of thumb that is not conditioned on market state is a far more useful rule than one that is.

We sort the 3512 windows into quintiles, first by  $\bar{\lambda}$  and then by the fitted  $n$ , and report the corpus median of the normalised single-stage tail  $p_{99}/T$  within each quintile. Normalising by  $T$  removes the service floor, so a value of 1.0 means no tail at all and the quintile-to-quintile ratio isolates the conditioning effect.

|                                                          |        | $p_{99}/T$ , corpus median |             |             |             |             |
|----------------------------------------------------------|--------|----------------------------|-------------|-------------|-------------|-------------|
| quintile                                                 | median | $T=8$                      | $T=16$      | $T=32$      | $T=64$      | $T=128$     |
| <i>by packet rate <math>\bar{\lambda}</math> (pkt/s)</i> |        |                            |             |             |             |             |
| Q1                                                       | 202    | 1.08                       | 2.29        | 4.74        | 12.25       | 31.76       |
| Q3                                                       | 451    | 1.07                       | 2.12        | 4.48        | 12.68       | 42.81       |
| Q5                                                       | 942    | 1.08                       | 2.11        | 4.51        | 13.92       | 55.67       |
| <b>Q5/Q1</b>                                             |        | <b>1.00</b>                | <b>0.92</b> | <b>0.95</b> | <b>1.14</b> | <b>1.75</b> |
| <i>by branching ratio <math>n</math></i>                 |        |                            |             |             |             |             |
| Q1                                                       | 0.769  | 1.08                       | 2.18        | 4.51        | 11.61       | 30.73       |
| Q3                                                       | 0.798  | 1.08                       | 2.16        | 4.61        | 13.17       | 42.83       |
| Q5                                                       | 0.825  | 1.07                       | 2.15        | 4.61        | 14.65       | 58.73       |
| <b>Q5/Q1</b>                                             |        | <b>1.00</b>                | <b>0.98</b> | <b>1.02</b> | <b>1.26</b> | <b>1.91</b> |

The result splits cleanly at  $T \approx 32\text{--}64\mu\text{s}$ . For  $T \leq 32\mu\text{s}$  the normalised tail is flat in both variables: a  $4.66\times$  change in packet rate moves it by 0–8%, and in the wrong direction at  $T = 16$  and  $T = 32$ , where the *busiest* quintile has the slightly *smaller* normalised tail. A  $0.769 \rightarrow 0.825$  change in branching ratio moves it by at most 2%. For  $T \geq 64\mu\text{s}$  both variables bite, and in the expected direction: by  $T = 128\mu\text{s}$  the busiest quintile’s tail is  $1.75\times$  the quietest and the most critical quintile’s is  $1.91\times$  the least. This is the same boundary that separates the clustering regime from the utilisation regime elsewhere in this paper, and it should be: for  $T \geq 64\mu\text{s}$  the Poisson null has a tail too (Section 4.1), utilisation  $\rho = \bar{\lambda}T$  is the governing quantity, and  $\rho$  is proportional to  $\bar{\lambda}$  by construction. Intensity conditioning matters exactly where classical queueing theory says it should, and nowhere earlier.

**Why intensity is inert below the crossover.** The null result is not a measurement artefact of the quintile split; it has a direct mechanical explanation in the arrival stream. Taking the same  $\bar{\lambda}$  quintiles and reporting the corpus median of the per-window interarrival-gap quantiles:

| quintile     | $\bar{\lambda}$ (pkt/s) | $p_1$ gap     | $p_5$ gap     | $p_{25}$ gap  | $p_{50}$ gap  | mean gap      |
|--------------|-------------------------|---------------|---------------|---------------|---------------|---------------|
| Q1           | 202                     | 7.45          | 8.85          | 21.30         | 87.10         | 4944          |
| Q2           | 320                     | 7.52          | 8.75          | 21.87         | 84.52         | 3128          |
| Q3           | 451                     | 7.53          | 8.70          | 22.70         | 82.16         | 2215          |
| Q4           | 620                     | 7.51          | 8.66          | 23.36         | 81.98         | 1612          |
| Q5           | 942                     | 7.43          | 8.58          | 22.55         | 76.52         | 1062          |
| <b>Q5/Q1</b> | <b>4.66×</b>            | <b>0.997×</b> | <b>0.969×</b> | <b>1.059×</b> | <b>0.878×</b> | <b>0.215×</b> |

All gaps in microseconds. The mean gap falls by exactly the reciprocal of the rate increase,  $0.215 \approx 1/4.66$ , which is definitional. Every other column is approximately invariant, and the tightest one is the most invariant of all: the  $p_1$  gap moves from 7.45 to  $7.43\mu\text{s}$ , a ratio of 0.997, while the rate more than quadruples.

This is what a Hawkes process does, and it is worth stating plainly because it is the crux of the whole regime. Raising the intensity of a self-exciting process raises the rate at which *clusters* are born; it does not repack the events *inside* a cluster. The intra-cluster spacing is set by the exchange’s packet-emission granularity — a property of the CME gateway, not of the market state — and on this corpus that floor sits at roughly  $7.5\mu\text{s}$ . A busy window is a window with more bursts in it, not a window with tighter bursts.

The queueing consequence follows immediately. A deterministic server of service time  $T$  begins to build a queue when it is handed two arrivals separated by less than  $T$ , so the service time at which a tail appears is set by the *tight* end of the gap distribution, not by its mean. The tight end is invariant to  $\bar{\lambda}$ ; therefore the threshold is invariant to  $\bar{\lambda}$ . And the two numbers agree: the measured  $p_1$  gap of  $7.43\text{--}7.53\mu\text{s}$  across every intensity quintile brackets the measured tail onset of Section 4.5, between  $T = 7$  and  $T = 8\mu\text{s}$ , where the fraction of windows carrying any tail jumps from 4.5% to 98.6%.

**Robustness: the null is not an artefact of the window length.** A 30-minute window averages  $\bar{\lambda}$  over a long stretch, and one could reasonably object that the quintile split above is comparing time-averaged rates that have smoothed away the intensity variation that matters. We therefore reran the entire pipeline — packet extraction, per-window Hawkes MLE, Poisson null and the full  $(T, N)$  grid — on *five*-minute windows, giving 21,060 windows in place of 3512. Shorter windows do resolve more intensity variation, as expected: the corpus  $\bar{\lambda}$  spread widens from  $6.8\times$  to  $7.3\times$  between the 5th and 95th percentiles, from  $26.0\times$  to  $30.6\times$  between the 1st and 99th, and the observed maximum rises from 2152 to 2692 packets per second; the 5th–95th range of the fitted branching ratio widens from 0.0717 to 0.0786. The conditioning result is nevertheless unchanged, to two decimal places, across the entire sweep:

|                                | $\bar{\lambda}$ Q5/Q1 | $T=8$ | $T=16$ | $T=32$ | $T=64$ | $T=128$ |
|--------------------------------|-----------------------|-------|--------|--------|--------|---------|
| 30-min windows ( $n = 3512$ )  | $4.66\times$          | 1.00  | 0.92   | 0.95   | 1.14   | 1.75    |
| 5-min windows ( $n = 21,060$ ) | $4.98\times$          | 1.00  | 0.93   | 0.96   | 1.15   | 1.74    |

Entries are the Q5/Q1 ratio of the normalised single-stage tail  $p_{99}/T$ . Resolving intensity six times more finely, over a wider realised range of it, moves no entry by more than 0.01. The inertness of intensity below the crossover is a property of the arrival process, not of the measurement window.

**Consequence for the design rule.** The rule of thumb this paper recommends — split a stage whose service time exceeds the feed’s tight-gap floor, leave it alone below — can therefore be stated as a constant of the feed rather than as a function of market state. It does not need a runtime intensity estimate, it does not change between a quiet morning and a volatile open, and a system tuned on one is tuned for the other. That is a stronger and simpler claim than the intensity-conditioned rule we set out to find, and it holds only because of the invariance measured above; it should be rechecked on any feed whose packet-emission behaviour differs from CME MDP3’s, since it is that emission granularity, and not anything about the order flow, that sets the constant.

Note also what the tight-gap floor is *not*: it is not the mean interarrival gap. On this corpus the mean gap is  $2215\mu\text{s}$  at the median window, some  $300\times$  the  $p_1$  gap. A rule that said “split stages longer than the average interarrival time” would recommend never splitting anything in an HFT hot path, and would be wrong by more than two orders of magnitude. The averages are irrelevant to the tail; only the tight end of the gap distribution is.

## 4.7 Sensitivity to $h$

The design equation is monotone in  $h$ : at fixed  $T$  and fixed corpus  $\Delta(N)$ , every  $\kappa$  breakpoint in the table above scales linearly with  $h$ , so halving  $h$  halves every breakpoint and moves  $N^*$  up one notch at every  $T$  for a fixed  $\kappa$ . At  $h = 1.7 \mu\text{s}$  the  $p_{99}$ -minimising  $N$  is 1, 1, 1, 2, 4, 8, 8 across the seven service floors; at a hypothetical  $h = 0.5 \mu\text{s}$  busy-poll mailbox it would be 1, 1, 2, 2, 4, 8, 8 — the change is a single notch, at  $T = 8 \mu\text{s}$ , because below the tight-gap floor no reduction in  $h$  can make a split worth taking when there is no tail to remove. Cheapening the hop widens the band of service times over which splitting pays; it does not move the threshold itself, which belongs to the feed. For any other framework the procedure is the same: measure the one-way asynchronous hop cost  $h$  on its own dispatch path, put it into Equation (11) with the corpus  $\Delta$  of Table 2, and read off  $N^*$ ; a framework whose  $h$  exceeds the tail excess a split could remove gets  $N^* = 1$ , and the recommendation reverses.

**$h$  is load-dependent;  $1.7 \mu\text{s}$  is a conservative upper bound.** The  $h = 1.7 \mu\text{s}$  figure is the one-way cost of the shipping KASPAR-HFT BQueue mailbox measured at low load, where the receiver is sleeping on a condition variable and the send path must pay the kernel wake latency to bring it back onto CPU. That wake latency is the dominant term in the  $\sim 3.4 \mu\text{s}$  round-trip of Table 4 in Maciejewski [32], and the general fact that a spinning receiver delivers a message an order of magnitude faster than a wake-from-condvar receiver is standard in the microsecond-latency systems literature: David et al. [58] measure the primitive costs directly on commodity Linux hardware, Belay et al. [59] exploit this in the IX data plane by pinning busy-polling network threads, and Rigtorp’s operational guide [37] formalises the “run one thread per core and busy-wait” rule that flows from it; the fast\_send paper’s own grouped-send row (Table 4,  $\sim 90 \text{ ns}$ ) is the same effect on the same framework — when the receiver never actually sleeps, the send path collapses. Under sustained load, and specifically inside a Hawkes cluster where the receiver’s queue is never empty, the receiver never sleeps: successive messages arrive faster than the receiver’s per-message service time, so it stays awake in a tight run-loop and the enqueue path reduces to a mutex acquire, a queue push, and a condition-variable notify against a receiver that is not waiting (a near-no-op). The effective  $h$  in that regime is dominated by the mutex + cache-line-transfer cost, typically a few hundred nanoseconds rather than  $1.7 \mu\text{s}$ . This is exactly the operating regime in which the paper’s tail-latency win is largest — the win compounds while  $h$  shrinks — so the paper’s reported effect underestimates the true benefit whenever the pipeline is inside a burst. Characterising the load-conditioned hop cost  $h(\rho)$  is left to future work; a full accounting would strengthen the results reported here.

## 5 Evaluation: do empirics match the analytical predictions?

This section compares the empirical numbers of Section 4 against the closed-form predictions of Section 2.

**Proposition 1 (burst-limit throughput).** The single-message identity  $p_{50}(N) = T + (N - 1)h$  is predicted exactly by Proposition 1 for any arrival law when the median message finds the queue empty. Table 1 confirms it in all 28  $(T, N)$  cells of the Poisson arm and in 25 of the 28 Hawkes cells at  $h = 1.7 \mu\text{s}$ , to the  $0.01 \mu\text{s}$  reporting precision; the three Hawkes exceptions are the longest-service cells, where the median message itself waits (Section 4.2).

**Theorem 1 (reduction and scaling bound).** The prediction is  $\Delta(N) \leq \Delta(1)/N$  on every window, and more strongly  $W_j^{(N)} \leq W_j^{(1)}/N + (N-1)h$  for every individual message. The corpus-median ratios in Table 2 sit at or below  $1/N$  in every cell, with substantial slack everywhere: the tightest ratio in the sweep is 0.163 at  $T = 32$ ,  $N = 2$ , against a bound of 0.5. The slack is largest where a split drives the per-stage service under the tight-gap floor and removes the excess outright rather than dividing it, which is why the fitted  $\gamma$  is 3.32 and 2.68 at  $T = 64$  and 128 rather than near 1. The pathwise form was checked directly by the simulator: across all 3512 windows, all seven service floors, both arms, and  $N \in \{2, 4, 8\}$ , no message violates the bound (the check allows  $N$  nanoseconds for the rounding of  $T/N$ ; the largest excess over that tolerance is 0). The reduction identity Equation (4) is visible in Table 2 without any fitting:  $\Delta$  is constant along every diagonal of constant  $T/N$ , so the tandem’s tail is the single-server tail at  $s_{\max}$  on real arrivals exactly as on synthetic ones (Appendices B and C).

**Corollary 2 (Poisson null).** The prediction is  $p_q^P(N) = T + (N-1)h$  wherever the per-stage utilisation satisfies  $\rho/N < 1 - q$ , and a visible wait wherever it does not. Table 1 matches both halves: Poisson  $p_{99}$  equals the floor in 22 of the 28 cells and exceeds it in exactly the six cells whose per-stage utilisation clears  $10^{-2}$  —  $(T, N) = (32, 1), (64, 1), (64, 2), (128, 1), (128, 2)$  and  $(128, 4)$ , at  $\rho/N = 0.0144$  to  $0.0578$  — and in no cell below it. Poisson  $p_{99.9}$  likewise exceeds the floor exactly in the cells with  $\rho/N > 10^{-3}$  (single stage at  $T \geq 4$ , two stages at  $T \geq 8$ , four stages at  $T \geq 16$ ) and nowhere else. The null arm is therefore a working queue whose small wait appears where PASTA says it must, and splitting under it costs  $(N-1)h$  at every reported quantile.

**Verdict.** All three analytical statements survive the full-corpus test. The quantitatively load-bearing one, the pathwise bound of Theorem 1, holds on every message of every window, and its slack is measured rather than assumed:  $\gamma = 3.32$  and  $2.68$  at  $T = 64$  and  $128 \mu\text{s}$ , and at  $T \leq 32 \mu\text{s}$  the tandem removes the excess outright rather than dividing it.

## 6 Cross-validation: the same arrival process, measured live

Everything in Part II so far is a measurement on *recorded* data: pcap-derived message tapes, replayed offline through the KASPAR-HFT simulator. That design buys scale and repeatability, and it costs external validity. A replayed tape cannot show whether the clustering we fit survives the network path to a receiver, whether it is a property of the 2025–2026 sample we happened to hold, or whether the service times the sweep assumes bear any relation to what a production decoder actually spends.

We are able to answer all three, because we also instrumented the production system. This section compares the two measurements. They were built for different purposes and share almost nothing beyond the exchange they observe, which is what makes the agreement between them worth reporting.

### 6.1 Two measurements, independent where it counts

The live measurement instruments the KASPAR-HFT market-data path in production: a timestamp  $t_0$  when a UDP datagram is handed up from the socket, a timestamp  $t_1$  when the resulting book update is published, and one 16-byte record per admitted message carrying the pair together with the ring-buffer depth and the message’s position inside its datagram. It covers 8.29 million messages on six streams — ES, NQ and ZN futures, book and trade — over 53 minutes of a single afternoon. Its latency results are reported in Maciejewski [32]; we draw on them here, and on the underlying per-message records, as a control on Part II.

|               | Part II corpus                       | production measurement              |
|---------------|--------------------------------------|-------------------------------------|
| source        | pcap tapes, replayed offline         | live multicast, production recorder |
| period        | 2025-01 to 2026-02                   | 2026-09-16 (7 months out of sample) |
| instruments   | NQ front month                       | ES, NQ, ZN $\times$ book, trade     |
| arrival clock | CME gateway <code>sendingTime</code> | receiver $t_0$ at socket read       |
| estimator     | per-window, median of 3512 windows   | pooled over one 53-min capture      |
| sample        | 3512 windows, 276 sessions           | 8.29 M messages, 6 streams          |

**Table 3.** The two measurements compared in this section. They share the exchange and nothing else: different instruments, a period that does not overlap, a different arrival clock, and a different estimator.

The arrival-clock row carries the most weight. Part II keys arrivals on `sendingTime`, the timestamp the CME gateway writes into the UDP header as the packet leaves the exchange. The live measurement keys them on  $t_0$ , taken by our own process after the packet has crossed the network, the NIC and the kernel. If the clustering reported in Section 3.2 were an artefact of exchange-side emission scheduling, the intervening path — with its own queueing, coalescing and interrupt moderation — would be expected to smooth it. It does not.

### 6.2 What agrees

Table 4 puts the four statistics that both measurements compute side by side. The NQ column is the like-for-like comparison, since Part II’s corpus is NQ front month; the third column gives the range across the three live book streams.

|                                  | Part II (NQ, 3512 win) | live NQ book                              | live, three books |
|----------------------------------|------------------------|-------------------------------------------|-------------------|
| $P(\text{gap} < \text{mean}/10)$ | 73.2%                  | 63.1%                                     | 63.1–84.8%        |
| Fano(5s) collapse under shuffle  | 14.8×                  | 12.2×                                     | 12.2–18.5×        |
| Hurst exponent $H$               | 0.740                  | 0.702                                     | 0.702–0.743       |
| branching ratio $n$              | 0.798 (MLE)            | $\leq 0.850\text{--}0.972$ (Fano ceiling) |                   |

**Table 4.** The same four Poisson-rejection statistics, computed independently on recorded and on live data. The Poisson values are 9.52%, 1×, 0.5 and 0 respectively.

Three of the four are direct comparisons and agree to within the spread across live streams. The fourth is not a comparison of like with like and is the more informative for it. Part II fits the branching ratio by maximum likelihood, per window, using the Ogata recursion (Section 3.2); the live measurement never fits a Hawkes process at all, and instead bounds  $n$  from above by inverting the Fano ceiling,  $n \leq 1 - 1/\sqrt{F}$ . The two estimators share no machinery. The MLE’s 0.798 sits inside every bound the live streams imply, which run from 0.850 on NQ trade to 0.972 on NQ book. The bound is loose — it is an inequality derived from a second moment — and the agreement is therefore weak evidence taken alone; its value is that it is evidence of a completely different kind, and it does not contradict.

What the section as a whole establishes is that the arrival process Part II models is not an artefact of replay, of the sample period, or of the exchange-side clock. It is present, at the same strength, in a live receiver seven months later on instruments the corpus never contained.

### 6.3 The service floor lands on the switch-on threshold

The live measurement supplies one quantity Part II cannot obtain from tapes at all: what a real decode of a real message actually costs. Isolating messages that arrive to an empty ring and first in their own datagram — so that neither queueing nor in-packet serialisation contributes — gives the hot path directly. Three estimators constructed differently agree:

| method                                                 | ES   | NQ   | ZN   |
|--------------------------------------------------------|------|------|------|
| direct median, empty ring and first in packet          | 7.01 | 7.24 | 6.89 |
| fitted intercept of the median position ladder         | 6.98 | 7.23 | 6.83 |
| floor at single-message packets during an FOMC release | 7.1  | 7.5  | 7.0  |

**Table 5.** Three independent estimates of the KASPAR-HFT socket-to-book floor, in  $\mu\text{s}$ , on the book streams. The three differ in sample, estimator and load regime. The third was taken while the packet rate rose 4.6 to 9.8× within a second; the floor moved by at most +0.4  $\mu\text{s}$  and fell on two of the three instruments.

Across all six streams the fitted floor spans 6.81 to 7.23  $\mu\text{s}$ , a range of 0.42  $\mu\text{s}$  from six separate multicast channels sharing nothing but the binary.

Now place that against Section 4.6. The service time at which a deterministic stage on this feed begins to accumulate a tail is the feed’s tight-gap floor, its 1st-percentile interarrival gap, measured at 7.43 to 7.53  $\mu\text{s}$  and invariant across intensity quintiles; the refined sweep of Section 4.5 locates the switch-on between  $T = 7$  and  $T = 8 \mu\text{s}$ . The production decoder’s floor is 6.81–7.23  $\mu\text{s}$ .

*The system sits on the threshold.* Not near it as an order-of-magnitude coincidence — on it, to within a few hundred nanoseconds, by two measurements that have no common term.

It is worth being precise about where the crossing happens, because the live measurement is sharp enough to locate it. Per-packet service is the floor plus one slope for each message beyond the first

(Section 6.5), so the service time of a datagram carrying span messages is  $\text{floor} + (\text{span} - 1) \times \text{slope}$ . Evaluating that on each stream’s own fitted pair:

|          | span 1    | span 2    | span 3    |
|----------|-----------|-----------|-----------|
| ES book  | 6.98      | 7.55      | 8.11      |
| NQ book  | 7.23      | 7.54      | 7.85      |
| ZN book  | 6.83      | 7.80      | 8.76      |
| ES trade | 6.81      | 7.52      | 8.24      |
| NQ trade | 7.14      | 7.67      | 8.19      |
| ZN trade | 6.81      | 7.77      | 8.74      |
| range    | 6.81–7.23 | 7.52–7.80 | 7.85–8.76 |

**Table 6.** Per-packet service time in  $\mu\text{s}$  by datagram span, from each stream’s measured floor and slope, against the feed’s tight-gap floor of 7.43–7.53  $\mu\text{s}$  (Section 4.6). Every stream is below the threshold at span 1 and above it at span 2.

All six streams are below the threshold when a datagram carries one message and above it when it carries two, with no stream crossing anywhere else. That is six independent crossings landing in the same unit interval of span, from floors and slopes fitted separately on six multicast channels and a threshold measured on a different corpus entirely. The design rule of Section 10 is therefore not a hypothetical for this system: the operating point is the boundary condition, and which side of it a given message falls on is decided by whether it arrived alone in its packet — which is decided, in turn, by whether it arrived inside a cluster.

## 6.4 Where the two disagree, and what each correction is

Two quantities do not match, and in both cases the discrepancy is a measurement artefact that one of the two measurements is positioned to diagnose. We report both rather than the agreements alone.

**The absolute Fano factor: pooling inflates it.** Part II measures a corpus-median Fano at a 5-second window of 259.7; the live streams give 1024 to 1260. The ratio under the gap shuffle agrees closely (Table 4), so the discrepancy is in the level, not in the effect. The cause is the estimator row of Table 3. The live figure is computed over one uncut 53-minute capture that spans a quiet stretch and a busy one; a rate that drifts between the two inflates the count variance and reads as clustering whether or not the process is self-exciting. Part II computes the Fano factor *within* each 30-minute window and then takes the corpus median, which removes between-window rate variation by construction. Pooling costs roughly a factor of four here, and the windowed estimate is the conservative one; the clustering result does not depend on which is used, but the level does, and 259.7 is the number we stand behind.

**The shuffled Hurst null: finite-sample bias, and Part II bounds it.** A Fisher–Yates shuffle of a gap sequence destroys the ordering while preserving the marginal exactly, so the shuffled series is a renewal process and its Hurst exponent must be 0.5. The live measurement recovers 0.600 to 0.660 instead, and attributes the excess to finite-sample bias in the Hurst estimator at the number of windows available in a 53-minute capture — which is why it states its result as the gap between 0.74 and 0.62 rather than as the distance from 0.5. That attribution was a reasonable inference rather than a demonstration. Part II can demonstrate it: over 3512 windows the same shuffle returns

$H = 0.501$ , and the uniform-shuffle Poisson null returns  $H = 0.496$  (Section 3.2). The bias is real, it is finite-sample, and it washes out with more and shorter series. The live study’s caveat was correct and its headline gap is the right way to read it.

## 6.5 What the live data contradicts, and what survives

One assumption of Part II’s simulator does not survive contact with the production measurement, and it should be stated plainly.

The sweep of Section 3.5 treats each packet as one arrival requiring a *constant* service time  $T$ . The live data shows that per-packet service is nothing of the kind. A UDP datagram on this feed carries a variable number of SBE messages — call it the packet’s *span* — which are decoded in order, so a message at position  $k$  waits for  $k$  decodes before its own. Median latency follows

$$\text{median} = \text{floor} + \text{slope} \times \text{position},$$

with the floor of Section 6.3 and a slope of  $0.31$  to  $0.97 \mu\text{s}$  per message depending on the stream. Span is not close to constant:

| stream   | packets   | messages  | mean | $p_{50}$ | $p_{90}$ | $p_{99}$ | max |
|----------|-----------|-----------|------|----------|----------|----------|-----|
| ES book  | 2,123,042 | 2,415,610 | 1.19 | 1        | 1        | 7        | 45  |
| NQ book  | 2,506,205 | 3,010,036 | 1.59 | 1        | 2        | 11       | 36  |
| ZN book  | 824,536   | 965,334   | 1.25 | 1        | 1        | 6        | 45  |
| ES trade | 83,140    | 235,601   | 4.10 | 3        | 8        | 17       | 83  |
| NQ trade | 50,279    | 83,233    | 3.47 | 2        | 7        | 16       | 85  |
| ZN trade | 22,864    | 100,659   | 5.86 | 2        | 11       | 71       | 85  |

**Table 7.** Messages per UDP datagram. The median book packet carries one message and the 99th percentile carries 6 to 11; on ZN trade the 99th percentile carries 71. A mean of 1.19 describes neither end of that distribution.

Because service is charged per message, the large packets carry a share of total processing time far out of proportion to their number:

| span $\geq 10$ | % of packets | % of messages | % of decode time |
|----------------|--------------|---------------|------------------|
| ES book        | 0.52         | 4.55          | 5.34             |
| NQ book        | 1.68         | 4.87          | 2.57             |
| ZN book        | 0.47         | 7.29          | 7.29             |
| NQ trade       | 5.34         | 19.31         | 16.16            |
| ES trade       | 6.69         | 26.53         | 40.10            |
| ZN trade       | 11.31        | 66.18         | 85.54            |

**Table 8.** Concentration of work in large datagrams. On ZN trade, 11.3% of packets carry 66% of the messages and consume 86% of decode time.

**Remark** (Reconciliation pending on Table 8). *The message counts in Table 7 do not equal the product of the packet count and the stated mean span, and they run 12–24% below the per-stream message counts used elsewhere in the live measurement, which suggests a difference in admission rule or weighting between the packet-level and message-level accounting that we have not yet traced. The qualitative claim — that decode time concentrates in a small minority of large datagrams, overwhelmingly so on the trade streams — is robust to any resolution of this, since it is visible*

directly in the span quantiles. The precise shares in Table 8 should be treated as provisional until the accounting is reconciled. No argument elsewhere in this paper depends on them.

So the deterministic-service assumption is violated at the packet level, and Part II’s  $T$  should be read as the per-packet work at a corpus-typical span rather than as a constant of the system. Two things nonetheless survive, and the second is a strengthening rather than a repair.

First, at the *message* level the assumption is close to right: the slope is fixed within a stream to the precision of the fit, and it is the composition of packets, not randomness in per-message work, that makes per-packet service vary. Theorem 1 is pathwise and arrival-law independent, and Remark 2.2 already handles unequal stages through  $s_{\max}$ ; what the live data adds is a measured magnitude for the variability that Section 2.3 lists as an assumption.

Second, span is a genuinely separate tail mechanism operating on the same stream, and it is active precisely where the queueing mechanism is not. Rather than leave that as a caveat we removed the assumption and re-ran the whole sweep with span-dependent service; Section 6.6 reports the comparison. The short version is that the serialisation channel supplies the entire tail below the tight-gap floor, where clustering supplies none, and is negligible above it, where clustering supplies all of it — and that the two turn out to be close to independent on this corpus rather than compounding, contrary to what we expected.

## 6.6 Span-aware service: re-running the sweep with the assumption removed

The objection above is testable rather than merely acknowledgeable, so we tested it. We re-extracted the per-packet span for all 3512 windows — aligned element-wise against the cached arrival arrays, with every window’s span count checked against the packet count already in the cache — and re-ran the entire sweep with service that depends on it.

Three service models are compared. All three hold the *mean* service per window at  $T$ , so the utilisation  $\rho$  is identical across them and only the distribution of work across packets differs; any difference in the tail is the span effect and not a slower system.

|         |                                               |                          |
|---------|-----------------------------------------------|--------------------------|
| const : | $S_i = T$                                     | (the Part II assumption) |
| span :  | $S_i = A(1 + r(\sigma_i - 1))$ , $r = 0.0432$ | (measured slope/floor)   |
| prop :  | $S_i = T \sigma_i / \bar{\sigma}$             | (upper bound)            |

with  $A$  set so that the window mean of  $S_i$  is  $T$ . The *span* model uses the ratio of slope to floor measured on NQ book in Section 6.3, 0.312/7.23. The *prop* model removes the per-packet fixed cost entirely, making service proportional to the message count; that is physically too strong, since the measured floor dominates the slope by more than twenty to one, and it is carried only to bracket the effect from above. As a correctness check the *const* arm reproduces the published constant-service grid of Table 1 to zero difference in all 28 cells.

**The unit matters more than the model.** On this corpus 96.1% of packets carry exactly one message and the mean span is 1.073, so a *packet*-weighted quantile is almost entirely a statement about singletons and barely moves under any service model. But packets are not what a strategy trades on. Weighting by message — which counts a 20-message datagram twenty times, as the decoder does — the in-packet position distribution is concentrated at zero and has a long upper tail:

| message quantile   | $p_{50}$ | $p_{90}$ | $p_{95}$ | $p_{99}$ | $p_{99.9}$ | max |
|--------------------|----------|----------|----------|----------|------------|-----|
| in-packet position | 0        | 0        | 1        | 4        | 19         | 84  |

93.2% of messages are alone in their packet and pay no serialisation at all; the remaining 6.8% constitute the entire tail. This is the same shape the live measurement reports, and it is why the comparison below is stated message-weighted.

| $T$ | const    |          |            | span     |          |            | span/const |            | prop/const |            |
|-----|----------|----------|------------|----------|----------|------------|------------|------------|------------|------------|
|     | $p_{50}$ | $p_{99}$ | $p_{99.9}$ | $p_{50}$ | $p_{99}$ | $p_{99.9}$ | $p_{99}$   | $p_{99.9}$ | $p_{99}$   | $p_{99.9}$ |
| 2   | 2.00     | 2.00     | 2.00       | 1.99     | 2.34     | 3.54       | 1.17       | 1.77       | 4.64       | 23.35      |
| 4   | 4.00     | 4.00     | 4.00       | 3.99     | 4.67     | 7.08       | 1.17       | 1.77       | 4.69       | 26.73      |
| 8   | 8.00     | 8.56     | 10.10      | 7.97     | 9.63     | 14.87      | 1.12       | 1.47       | 5.21       | 24.64      |
| 16  | 16.00    | 34.00    | 53.55      | 15.95    | 35.08    | 56.62      | 1.03       | 1.06       | 3.36       | 11.53      |
| 32  | 32.00    | 143.74   | 256.51     | 31.89    | 144.41   | 260.14     | 1.00       | 1.01       | 2.37       | 5.88       |
| 64  | 98.80    | 815.22   | 1557.72    | 100.02   | 820.05   | 1573.97    | 1.01       | 1.01       | 2.37       | 3.94       |
| 128 | 366.78   | 5440.22  | 12844.51   | 371.91   | 5550.42  | 13159.71   | 1.02       | 1.02       | 2.51       | 3.10       |

**Table 9.** Message-weighted single-stage latency in  $\mu\text{s}$ , corpus median over 3512 windows, Hawkes arrivals, under constant and span-aware service at equal mean service. The last two columns give the *prop* upper bound as a ratio to const.

**The two tail mechanisms occupy disjoint regimes.** Read the span/const ratio column down the table. It is largest exactly where the queueing tail is absent and smallest exactly where the queueing tail dominates. At  $T = 2$  and  $4\mu\text{s}$  the constant-service model has *no tail whatsoever* —  $p_{50} = p_{99} = p_{99.9} = T$ , the result reported in Section 4.1 — while the span-aware model has  $p_{99} = 1.17T$  and  $p_{99.9} = 1.77T$ . Below the feed’s tight-gap floor there is still a tail; it is simply not a queueing tail. By  $T \geq 32\mu\text{s}$ , where cluster queueing produces a  $p_{99}$  of  $4.5T$  and rising, span contributes 0–2% on top.

**At the production operating point the two are comparable, and span is the larger.** The table above is a sweep over hypothetical service times, and reading the  $T \geq 32\mu\text{s}$  rows as the headline would be misleading, because the live system’s measured floor is  $6.81\text{--}7.23\mu\text{s}$  (Section 6.3). The comparable row is  $T = 8\mu\text{s}$ . Decomposing the message-weighted  $p_{99}$  excess over the floor into the part the constant-service model already produced (queueing) and the remainder introduced by span-aware service (serialisation):

| $T$ | floor  | queueing excess | serialisation excess | serialisation share |
|-----|--------|-----------------|----------------------|---------------------|
| 2   | 1.99   | 0.00            | 0.34                 | 100%                |
| 4   | 3.99   | 0.00            | 0.69                 | 100%                |
| 8   | 7.97   | 0.56            | 1.09                 | 66%                 |
| 16  | 15.95  | 18.00           | 1.14                 | 6%                  |
| 32  | 31.89  | 111.74          | 0.78                 | 1%                  |
| 64  | 100.02 | 716.42          | 3.61                 | 1%                  |
| 128 | 371.91 | 5073.43         | 105.07               | 2%                  |

At  $T = 8\mu\text{s}$  serialisation supplies 66% of the  $p_{99}$  excess and queueing 34%. That is the same ordering the live measurement reports independently: evaluating the position ladder at each stream’s own  $p_{99}$  position accounts for 71% to 80% of the observed  $p_{99}$  on four of its six streams, with the remainder left to queueing and to stalls. Applying that definition to the simulator at  $T = 8$  gives 97%, higher than the live figure because the only residual the model contains is queueing — it has

no stall population. The agreement to take from this is the attribution, not the third digit: at the service time an HFT decode path actually runs at, in-packet serialisation is the larger of the two tail mechanisms, and both measurements say so.

This completes rather than contradicts Part II. The claim that survives is the one about *queueing*: below roughly  $7.5\ \mu\text{s}$  of service, arrival clustering does not build a queue, and splitting the servicing chain buys nothing. What the span-aware run adds is that a second, unrelated tail mechanism is active in that same regime — in-packet serialisation — which is deterministic, is not a queue, and cannot be removed by pipelining the arrival stream. Indeed  $N^*$ , the  $p_{99}$ -minimising stage count, is *unchanged* in every one of the seven service floors when service is made span-aware: 1, 1, 1, 2, 4, 8, 8 under both models. The design rule of Section 10 is unaffected.

**The in-packet ladder.** Holding the queue empty and walking the position inside the packet reproduces the live measurement’s linear ladder. At  $T = 8\ \mu\text{s}$  the corpus-median latency by position fits

$$\text{median} = 7.98 + 0.344 \times \text{position} \quad (\mu\text{s}),$$

recovering the  $0.345\ \mu\text{s}$  per-message cost put into the model to three significant figures and confirming that queueing contributes nothing at the median at this service time. The live NQ-book ladder is  $7.23 + 0.312 \times \text{position}$ ; the shapes agree and the levels differ by the difference between a  $7\ \mu\text{s}$  measured decode path and an  $8\ \mu\text{s}$  nominal service floor.

**The two channels are additive, not coupled — correcting an earlier claim.** It is natural to suppose, and an earlier draft of this section asserted, that span is itself produced by the clustering: arrivals bunch, the gateway coalesces more messages into one datagram, and the unlucky message is both queued behind a burst and buried inside a large packet. If that held, the two mechanisms would compound and the constant-service estimates of Part II would be conservative. On this corpus it does not hold.

Two tests. First, a control arm in which each window’s span sequence is randomly permuted across packets, preserving the span marginal and the arrival times exactly and destroying only the association between them. If large packets arrived preferentially inside clusters, breaking that association would lower the tail. It does not: the ratio of the real to the span-shuffled  $p_{99}$  is 1.000 at every  $T \leq 8\ \mu\text{s}$  and lies between 0.99 and 1.02 at every larger  $T$ , with no systematic sign. Second, measuring the association directly, the correlation between a packet’s span and the logarithm of the gap that preceded it has corpus median  $-0.014$ ; packets arriving after the tightest decile of gaps have mean span 1.010 against 1.043 for those arriving after the loosest decile, a ratio of 0.97. If anything the association runs weakly the other way.

The live measurement reports the same instrument dependence and should have been read more carefully: mean span rises steeply with queue depth on ZN, where a quiet book is filled only by genuine bursts, but *falls* with queue depth on ES and NQ, whose rings back up with many small packets. Part II’s corpus is NQ, and it behaves like the live NQ streams. The conclusion is therefore that the queueing channel and the serialisation channel are close to independent here, that they add rather than multiply, and that Part II’s constant-service tables are neither conservative nor optimistic on that account — they simply measure one of the two. Whether the channels couple on an instrument that behaves like ZN is open, and is the case worth checking before transporting these numbers to another feed.

## Part III

# Best practices: implications for HFT systems design

## 7 The reference instance: KASPAR-HFT and its `fast_send` pipeline

The recommendations in this Part are stated for HFT systems in general, but they are worked out on one concrete system, and the reader should know what that system is. KASPAR-HFT [32, 33] is an open-source C++20 trading system for CME futures built on a shared-nothing actor framework, available at <https://github.com/vincent212/kaspar-hft>: every component (socket reader, feed arbitrator, SBE decoder, order book, strategy, order manager, iLink session) is an actor with private state and a mailbox, and components interact only by passing messages. It is at once a production system (MDP3 multicast in, full order books reconstructed order by order, authenticated iLink 3 sessions out) and, with the same actor code driven from recorded packet captures, a deterministic replay simulator in its own right. Part II does not run on that simulator: its measurements are made on a standalone tandem-Lindley model of the pipeline (Section 3.3), driven by packet-arrival timestamps extracted from the same class of capture. The framework supplies the architecture under study and the calibration of  $h$ ; it is not in the measurement loop. The `fast_send` paper documents the four extensions that make the actor model usable at HFT latencies: `fast_send` itself, a synchronous delivery in which the sending thread runs the receiver’s handler inline and takes the reply as a value; actor `Groups`, which co-schedule a set of actors on one thread behind a single mailbox; mailbox queue implementations selectable per actor; and a memory pool that removes per-message allocation and its tail. The defining property of `fast_send` is *receiver transparency*: a handler is written once and has no way to tell whether it was invoked synchronously or asynchronously, on which thread, or whether its sender is blocked. That property is what Section 8 builds on.

The reference pipeline of that paper is the design context this paper revisits. It has three sections joined by exactly two cross-thread hops. Two socket-reader actors, one per redundant multicast channel, do nothing but drain their UDP sockets and hand datagrams to a buffer actor by an asynchronous `send`; that hop is asynchronous because a reader that blocked on its callee would drop packets under a burst. From the buffer, arbitration between the A and B feeds, SBE decode, order-book construction, and signal generation run as a single chain of `fast_sends` on one thread, with no queue and no context switch between stages. When the signal decides to act it hands the order to an order-management actor by a second asynchronous `send`, because the iLink write takes microseconds and blocking the decode thread on it would stall the market-data queue. The measured costs behind that layout are Table 4 of the `fast_send` paper: a cross-thread `send` round trip of  $\sim 3.4\ \mu\text{s}$  (a mailbox wakeup across cores), a same-thread grouped `send` of 90 ns, and a `fast_send` of  $\sim 30$  ns. With a servicing chain of  $\sim 7\ \mu\text{s}$  measured tick-to-book on live ES, NQ, and ZN feeds, the framework’s own per-message contribution is under 1% of the floor, and the same measurement already found that the latency tail was set by the arrival process, which was strongly clustered, with a Fano-implied branching ratio of 0.85–0.97. That observation is the seed of the present paper.

Read against Part II, the reference pipeline is the single-thread servicing chain of Section 1.2 realised with actors: the whole  $T$  sits on one thread between two hops that were accepted because the alternative, blocking the chain, was worse. The two existing hops show that the design was never opposed to asynchronous handoffs as such; it placed them where the cost of not having them was obvious. What this paper adds is that the same logic applies inside the chain, once the chain is long enough for the feed to queue behind it. The recommendations that follow are therefore stated

as changes to this particular pipeline: where to cut it, in what order to place the stages, which boundaries to leave on `fast_send`, and how to decide each of those by measurement rather than by reasoning at a single call site.

**The rule applied to this pipeline, as measured.** The design equation can be run on the reference chain directly, using the diagonals of Table 2 as a lookup of the single-server tail excess  $\Delta_{\text{single}}(s)$  against per-stage service  $s$ . Take the measured  $\sim 7\mu\text{s}$  socket-to-book chain and the natural boundaries inside it (arbitration, SBE decode, book application, signal), and suppose for the sake of the arithmetic that decode and book dominate at roughly  $3\mu\text{s}$  each with arbitration and signal near  $0.5\mu\text{s}$ . Then every partition available at those cut points has  $s_{\text{max}} \leq 7\mu\text{s}$ , and Part II says  $\Delta_{\text{single}}$  is zero there: 0.00 at  $s \leq 4\mu\text{s}$ , and at  $s = 7\mu\text{s}$  a tail on only 4.5% of windows with a median excess of 0.00 (Section 4.5). The single-thread chain has no queueing excess to remove, every cut returns exactly its hop tax, and  $N^* = 1$ .

That is not a disappointment; it is Section 6.3 read from the other side. The production chain sits on the feed’s tight-gap floor, and the paper’s own rule says a stage on the floor is not to be split. What the live measurement adds (Section 6.6) is that the tail this chain does carry at its operating point is mostly in-packet serialisation, which pipelining the arrival stream cannot touch, and which the fast-code craft of Section 10 can.

**The rule applied to this pipeline, once it grows.** The rule engages as soon as the chain is longer than the floor, which happens the moment a strategy stage of any substance is placed on it, or on a feed with a tighter gap floor. Take a  $16\mu\text{s}$  chain with the same four boundaries, decode and book now at  $6\mu\text{s}$  each and the two small stages at  $2\mu\text{s}$ :

| partition of the $16\mu\text{s}$ chain | $s_{\text{max}}$ ( $\mu\text{s}$ ) | hops | $\Delta_{\text{single}}(s_{\text{max}})$ ( $\mu\text{s}$ ) | $p_{99}$ excess + hop tax |
|----------------------------------------|------------------------------------|------|------------------------------------------------------------|---------------------------|
| single thread (16)                     | 16                                 | 0    | 18.59                                                      | 18.59                     |
| 8 + 8 (arb+decode   book+signal)       | 8                                  | 1    | 0.61                                                       | 2.31                      |
| 8 + 6 + 2 (three stages)               | 8                                  | 2    | 0.61                                                       | 4.01                      |
| 2 + 6 + 6 + 2 (all four split)         | 6                                  | 3    | 0.00                                                       | 5.10                      |

Every entry is read directly off Table 2:  $\Delta_{\text{single}}(16) = 18.59$ ,  $\Delta_{\text{single}}(8) = 0.61$ ,  $\Delta_{\text{single}}(6) = 0$ . Three things follow, and each is an instance of a general rule rather than a fact about this chain. The balanced two-way cut is decisive: it removes  $18\mu\text{s}$  of  $p_{99}$  excess for  $1.7\mu\text{s}$  of hop, an eightfold improvement in the tail objective. Going to three stages does not lower  $s_{\text{max}}$  at all — the  $8\mu\text{s}$  stage is still the bottleneck — so the second hop buys nothing and costs 1.7, which is Remark 2.2’s third consequence in concrete form. Splitting all four does lower  $s_{\text{max}}$  to 6 and removes the last  $0.61\mu\text{s}$  of excess, but pays 1.7 to do it, and loses. So the recommendation for a chain of this length is one cut, placed so that the two halves are as equal as the natural boundaries allow, with the larger half at ingress (Remark 2.2) — and the numbers above are a prior for what the end-of-build measurement of Section 8.3 should find, not a substitute for running it. The per-stage service times are illustrative; the operator’s own profile of the chain replaces them, and the lookup table does the rest.

## 8 Implementation on KASPAR-HFT: two dispatch paths, one handler

The central practical claim of this paper is that on Hawkes-clustered feeds the choice of `fast_send` versus asynchronous `send` at each actor boundary is not an early-stage architectural commitment but a late-stage optimization decision — one that should be made per boundary on the basis of

measured end-to-end median and tail latency of the assembled pipeline against the target feed, not on the basis of local reasoning at any single call site. KASPAR-HFT was designed to support this workflow. This section states the property that makes it possible, argues why it is critical from a system-optimization standpoint, and gives the operational discipline it enables.

The design equation maps onto framework choices in a specific way. The paper’s model of an  $N$ -stage tandem assumes each stage’s actor runs on its own thread and sends its output to the next stage through an asynchronous mailbox — that is what makes the pipeline pipelined, and the hop cost  $h$  is the wall-clock latency of that asynchronous handoff.

**Two hop paths in the shipping KASPAR-HFT framework.** Of the delivery mechanisms described in Section 7, two matter for the design equation. The default asynchronous `send` uses `BQueue`, a blocking mailbox built on `std::mutex + std::condition_variable + boost::circular_buffer`; a producer thread enqueues under the mutex and a consumer thread waits on the condition variable. This is the hop path the paper’s Section 3.3 model corresponds to, and halving its measured  $\sim 3.4\mu\text{s}$  round trip gives the one-way hop cost  $h \approx 1.7\mu\text{s}$  used throughout Part II.

The `fast_send` path is different: it is a *synchronous* dispatch that acquires a per-actor mutex, calls the receiver’s handler inline on the sender’s thread, and returns the reply. There is no queue and no cross-thread wake-up. Its per-call cost is much smaller than the `BQueue` path’s, but it is not a pipelined hop: the sender blocks until the receiver finishes, so a chain of `fast_sends` serialises through all stages on one thread and offers none of the throughput speedup of Proposition 1. The paper’s design equation does not apply to `fast_send` chains.

## 8.1 Receiver transparency

Every actor in KASPAR-HFT registers its handlers through the `MESSAGE_HANDLER(MsgType, handler_method)` macro. A handler is a function of one message argument; nothing in its body can access, observe, or condition on the delivery mode by which it was invoked. This property is called *receiver transparency* in Maciejewski [32] and is what makes the send-mode choice a purely sender-side decision.

Concretely:

- **Synchronous delivery:** `target->fast_send(m, this)`. The sender’s thread acquires `target`’s per-actor mutex, calls the receiver’s handler inline on the sender’s thread, and takes the reply as a return value. No queue, no cross-thread wake-up. Round-trip cost  $\sim 30\text{ ns}$  per Table 4 of Maciejewski [32].
- **Asynchronous delivery:** `target->send(m, this)`. The message is enqueued on `target`’s `BQueue` mailbox (`mutex + condition_variable + boost::circular_buffer`); `target`’s own thread wakes up on the condition variable and runs the handler. Round-trip cost  $\sim 3.4\mu\text{s}$  across threads, or  $\sim 90\text{ ns}$  same-thread within a `Group`, per the same reference.

Both paths deliver the same handler with the same message. Neither the handler’s body nor any actor downstream of it can tell which path fired. Switching a boundary between the two is a one-line change at the sender’s call site — `target->fast_send(m, this) ↔ target->send(m, this)` — with no code change on the receiver side, no header change, no interface change, no test change.

## 8.2 Why measurement, not local reasoning, must decide

The end-to-end latency profile of an HFT pipeline is a joint property of (i) which stages exist and where they are cut, (ii) which pairs of adjacent stages are joined by `fast_send` versus `send`,

and (iii) the arrival law of the input feed. Proposition 1 and Theorem 1 both fold across multiple boundaries: the throughput speedup is  $N$ -fold and the tail contraction is  $N^{-\gamma}$  with  $\gamma \geq 1$ , so the joint optimum is emphatically not deducible from local reasoning at any single boundary. A designer who reasons at one boundary at a time — "this hop is  $\sim 1.7 \mu s$ , that is a lot, use `fast_send`" — systematically underestimates the aggregate tail benefit of adding hops, and this is exactly the class of local reasoning that led the low-latency-systems community into the single-thread orthodoxy of Section 1.2.

Getting the joint optimum therefore requires end-to-end measurement of the assembled pipeline against a representative feed, and iteration over per-boundary flips. The observable objective is a two-tuple  $(p_{50}, p_{99})$  of end-to-end latency under the target feed, with an operator-chosen tradeoff between the two (Section 4.3's  $\kappa$ ). The design equation of Part II supplies a first-principles prior on where the optimum sits under a Hawkes feed, but the measurement is the arbiter, not the equation.

### 8.3 Operational discipline

The KASPAR-HFT design flow that this suggests, and which we now recommend as canonical for HFT actor pipelines running on self-exciting feeds, is:

1. Draft the pipeline as a chain of actors at the natural cut points of the servicing chain (`SocketReader`, SBE decode, order book, strategy, order manager on CME MDP3; four to five stages, with no artificial subdivision of any one of them).
2. Wire every intra-servicing-chain boundary with whichever of `fast_send` or `send` is convenient during development; do not commit to a topology at this stage.
3. Once the pipeline runs end-to-end on the target feed, measure the corpus-median  $p_{50}$  and  $p_{99}$  of end-to-end latency under the operator's live-representative tape, and record them as the baseline.
4. Perturb one boundary at a time — flip its `fast_send`  $\leftrightarrow$  `send` — and re-measure. Accept the flip if  $(p_{50}, p_{99})$  improves under the operator's chosen tradeoff.
5. Iterate until no single-boundary flip improves the objective.

Under the arrival law of a CME MDP3 feed the design equation of Section 4.3 predicts what steps 4–5 converge to: async `send` at every boundary whose removal would raise the largest stage, and `fast_send` everywhere else, in particular between short sub-tasks below the floor of Section 4.3 and between sub-tasks downstream of the bottleneck whose combined service stays at or below  $s_{\max}$  (Remark 2.2). The operator does not need to trust that prediction to run the discipline; the measurement is the arbiter, and the arrival-law prior tells the operator where to expect it to land.

### 8.4 Separation of code construction from latency optimization

A consequence of Section 8.3 that deserves separate emphasis: in KASPAR-HFT *two* distinct topology decisions are not knowable at design time. The first is the *actor-to-thread mapping* — which actors share a thread by being placed in a common `Group`, and which run on their own thread. The second is the *per-boundary send-mode choice* — whether each pair of adjacent actors is joined by synchronous `fast_send` or asynchronous `send`. Neither decision is deducible from the source code alone. Both are functions of the target feed's arrival statistics and the operator's median-versus-tail tradeoff, both of which are properties of the deployment environment. The optimum can only be measured from the assembled system running against a representative feed.

Under receiver transparency this is a feature rather than a limitation. Actor code, message-type headers, and the full suite of unit and integration tests can be written and run before either the actor-to-thread mapping or the send-mode choice has been decided. Code construction and testing therefore proceed independently of the latency-optimization consideration — both decisions are deferred to the end of the build, informed by measurements that could not have been taken any earlier because the system did not yet exist to measure. Correctness is established first; latency is optimized against the specific feed the system will run on, second. Two developers building two actors on the two sides of a boundary do not need to agree on whether their actors will ultimately live on the same thread or on different threads, nor on whether their handoff will be synchronous or asynchronous; they need to agree only on the message type. The message-type agreement is checkable at compile time; the two topology decisions are scalar dials adjusted at deployment.

This decoupling of correctness-phase engineering from latency-phase engineering is the practical payoff of receiver transparency, and it is what makes KASPAR-HFT’s design cycle qualitatively different from frameworks in which either topology decision leaks into the receiver’s interface (Section 8.5). Design proceeds without a commitment to a topology; the commitment is made once, near deployment, on the basis of measurements the code did not need to know about while it was being written.

## 8.5 Contrast with non-fungible frameworks

Most low-latency actor frameworks bake the sender’s choice of delivery mode into the receiver’s type. An Akka actor’s `tell` versus `ask`, a CAF actor’s `send` versus `request`, and an Erlang process’s `cast` versus `call` each expose different receiver-side surfaces, so flipping a boundary from sync to async is a small refactor rather than a one-line change. Under those frameworks the design flow of Section 8.3 is not tractable, and the operator’s incentive is therefore to lock in the send-mode choice early — reasoning locally — rather than to defer it to an end-of-build measurement pass. This structural incentive is a plausible contributing cause of the single-thread orthodoxy: frameworks that penalise deferral push designers toward the topology that is least sensitive to a wrong early guess, which is the collapsed single-thread servicing chain.

KASPAR-HFT’s joint mailbox API of `fast_send` and `send`, and the receiver-transparency property underneath, are the concrete machinery that make measurement-driven send-mode optimization practical. That is the framework property this paper’s design rule depends on operationally, and it is why we regard it as critical to system-level optimization on Hawkes-clustered feeds.

## 9 Other architectural considerations to re-evaluate under Hawkes

The paper’s central claim is that the single-thread hot-path recommendation, treated as an established convention in HFT infrastructure for over a decade, becomes wrong once the arrival law is recognised as Hawkes-clustered rather than Poisson (Section 1.2, Section 4). Two adjacent architectural conventions in the same low-latency-systems canon — pre-spinning receivers and pre-warming pipelines — rest on the same implicit Poisson-load assumption, and this section flags them as deserving the same treatment. We do not measure either intervention here; both belong to future work. What follows is an argument that the arrival law itself provides most of the benefit these interventions are engineered to deliver, on the specific quantile that HFT operators care about.

**Pre-spinning receivers.** The busy-poll recommendation of Rigtorp [37], Belay et al. [59], and the Aeron/Chronicle lineage (Section 1.2) pins a spinning thread per receiver so the wake path from

a sleeping condition variable is never taken. On a Poisson feed, any arrival can find the receiver asleep, and the wake latency is paid on average; pre-spinning eliminates it. On a Hawkes feed with branching ratio  $n \approx 0.8$ , sleep/wake transitions happen only at cluster boundaries. Every  $p_{99}$  message is intra-cluster (Section 4.1: the  $p_{99}$  message waits behind a cluster, since at these service times the Poisson null produces no wait at all), so by the time a tail-setting message arrives, the receiver has been continuously processing for many multiples of the per-message service time and is already effectively spinning as a natural consequence of the load. Under the model used here, pre-spinning therefore saves little or no wake latency at the tail; this is a prediction of the queueing model, not a measurement of a production receiver’s state, and should be evaluated directly. It saves wake latency for the median and for the minimum-latency floor (inter-cluster arrivals which do find the receiver asleep after seconds-scale gaps), but the saving is bounded by  $\sim 1.4 \mu\text{s}$  per hop (the  $1.7 \mu\text{s}$  wake-path cost less the few-hundred-nanosecond mutex and cache-line floor of Section 4.7) and is much smaller than the tail win from pipelining. There is also an operational cost: in the present author’s production experience, pinning spinning threads across the servicing chain has caused overall system degradation — steady core burn at 100%, cache thrash between the spin loop and the actor’s own working set, and pressure on shared resources like the LLC and memory bandwidth — that appears to eat back the modest median-latency gains, though this is anecdotal and not measured here.

**Pre-warming.** Pre-warming is arguably not a separate consideration but a subset of the wake-up cost already discussed above. The  $h = 1.7 \mu\text{s}$  one-way hop cost that this paper calibrates already captures the entire process-wake-up path on the receiver side: the kernel context switch, the reload of the receiver’s instruction cache, the L1/L2 misses on its actor state and message pool, the TLB refill, and the scheduler putting the thread back on a core. To the extent the thread hop is dominated by cluster-wait on Hawkes feeds — because every  $p_{99}$  arrival is intra-cluster and the receiver is, under the model, already awake — the cache-warmth components inside that  $h$  should follow the same pattern, since they are the same  $h$  decomposed. This is a consequence of the model rather than a measured result. Pre-warming interventions that fire synthetic traffic to keep instruction caches, data caches, TLBs, and branch predictors hot are therefore attacking the same  $h$  that Section 4 finds is dominated by cluster-wait. What such interventions can meaningfully deliver is confined to the median and the minimum — the inter-cluster first-arrival case where the receiver truly did cool down — and even there the delivery is bounded by the wake-up cost itself and comes at the price of continuous L1/L2 pollution from the warmup traffic. Under Hawkes the tail sees none of this.

**Why this needs measurement.** The convention of treating pre-spinning and pre-warming as first-order tail-latency techniques is, like the single-thread rule, a Poisson-load argument in disguise. On real self-exciting exchange feeds the arrival law itself provides the receiver-alive and cache-warm state on the quantile that matters (the tail), so both interventions are second-order at best on those quantiles and can be actively counterproductive once their steady costs are accounted for. Both claims are corollaries of the paper’s main result rather than separate contributions, but each deserves a dedicated empirical evaluation on Hawkes-clustered feeds, analogous to the single-thread evaluation this paper carries out. The two open questions are:

- What is the actual ROI of pinned busy-poll receivers on the Hawkes-load  $p_{50}, p_{99}, \text{max}$  triple, net of the steady core-burn cost and the cache pressure the spin loop imposes on the actor’s own working set?
- What is the actual ROI of periodic pre-warming traffic on the same triple, net of the L1/L2 pollution it creates and the residual advantage it delivers only on inter-cluster first-arrivals?

Both are direct extensions of the present paper’s methodology (measure under real CME MDP3 tapes; use a matched Poisson null to isolate the arrival-law dependence) and both are candidate follow-ons.

## 10 Recommendations for HFT system designers

The paper’s results overturn one architectural default (single-thread hot path) and cast serious doubt on two adjacent ones (busy-poll receivers, pre-warming). It is easy to read those results as an attack on the microsecond-shaving craft that HFT infrastructure engineers spend most of their time on, and they are not. This section separates what the paper leaves untouched from what it argues against, and states the design flow the paper’s design equation actually implies.

### 10.1 What the paper leaves untouched

The bit-level hot-path optimization craft — branch-free decoders, cache-line-aligned actor state, atomic-free single-writer queues, SBE decode in place, memory pools, `fast_send` at  $\sim 30$  ns, kernel-bypass NICs (Solarflare Onload, DPDK, VMA), `isolcpus`, `SCHED_FIFO`, huge pages, RTC counters instead of `clock_gettime`, non-blocking IO on the socket, avoidance of allocation on the hot path, careful use of the branch predictor — is every bit as necessary as it was before this paper. All of it shows up in the service floor  $T$ , which the paper takes as given and treats as the constant the multi-actor pipeline is compressing the tail *around*. The lower  $T$  is, the better the pipelining rule performs: Proposition 1 gives an  $N$ -factor throughput speedup at any  $T$ , and Theorem 1’s  $1/N$  bound applies to a single-stage wait  $\Delta(1)$  that itself scales with  $T$ . So the pipelining rule is complementary to, not a substitute for, the microsecond-shaving craft: a  $5\mu\text{s}$ -servicing-chain implementation of a two-stage tandem beats a  $10\mu\text{s}$ -servicing-chain implementation of the same two-stage tandem at every quantile, and the difference is entirely in the fast-code layer this paper does not touch.

### 10.2 What the paper does argue against

Two layers above the bit-level craft, both of which have been standard practice in the HFT community for a decade or more:

1. **Topology default: single-thread servicing chain.** The Poisson-orthodox recommendation to collapse the entire hot path onto one thread (Section 1.2). On Hawkes-clustered feeds this is empirically wrong: the tail is dominated by cluster-wait, and adding threads between actors along the servicing chain compresses that tail by  $N^{-\gamma}$  at a small  $(N - 1)h$  median cost, provided  $T$  sits above the floor of Section 4.3.
2. **Always-on interventions: pre-spinning receivers and pre-warming pipelines.** The Poisson-orthodox recommendation to burn a core-worth of CPU per receiver on busy-poll, and a smaller amount of CPU on periodic warmup traffic, to eliminate wake-up latency and cold-miss cost on the receiver’s side (Section 9). On Hawkes-clustered feeds the receiver is already effectively spinning and effectively warm at the arrival time of every  $p_{99}$  message, so these interventions target quantiles (median, minimum) that are not what an HFT operator cares about, and they pay a steady core-burn and cache-pollution cost for it.

These are orthodoxy about the *layer above* the fast-code craft, not about the fast-code craft itself. Practitioners who conflated “I made this loop 200 ns faster” with “therefore single-thread wins” were making two different claims at once; the paper only overturns the second.

### 10.3 Recommended design flow

The design flow that follows from the paper’s results, and that we recommend as the canonical HFT infrastructure workflow on self-exciting exchange feeds:

1. **Keep the fast-code craft.** Everything in the first subsection above. This is orthogonal to the paper’s recommendation and is still what pays the bills on absolute latency.
2. **Measure the target feed’s tight-gap floor, and use it as the split/don’t-split test.** Take the corpus of packet arrivals, form the per-window interarrival-gap distribution, and read off a low quantile of it — we use  $p_1$ . This single number is the service time at which a deterministic stage begins to accumulate a tail, and on CME MDP3 channel 318 it is  $\approx 7.5 \mu\text{s}$  (Section 4.6). A stage whose service time sits comfortably below it will not develop a queue and must not be split: the split buys nothing and costs  $h$  at every quantile. A stage above it should be split, subject to the tail excess exceeding  $h$ , which on this corpus happens about  $1.5 \mu\text{s}$  further along (Section 4.5). Two cautions. First, use the tight end of the gap distribution, never the mean: on this corpus the mean gap is  $2215 \mu\text{s}$ , some  $300\times$  larger, and a rule stated on it would say never to split anything. Second, this floor is a property of the exchange’s packet-emission granularity and is invariant to market intensity — it does not move between a quiet session and a volatile one (Section 4.6) — so it is measured once per feed, not tracked at runtime. Re-measure it for any feed whose gateway behaves differently.
3. **Draft the pipeline as a chain of actors at natural cut points, then balance it.** On a CME MDP3 handler: `SocketReader`, SBE decoder, order book, strategy actor, order manager — four or five stages, no artificial decomposition (Section 8). The tail factor is  $s_{\max}/T$  (Remark 2.2), so the goal of the decomposition is the smallest largest stage, and the largest stage belongs at ingress so that all queueing sits in one mailbox (Remark 2.2).
4. **Wire boundaries with whichever of `fast_send` or `send` is convenient during development.** Receiver transparency in KASPAR-HFT makes this a one-line-per-boundary decision that can be deferred (Section 8).
5. **Do not commit to an actor-to-thread mapping at design time.** Grouping decisions and per-boundary send-mode decisions are both scalar dials at deployment, not code-structure decisions (Section 8.4).
6. **Measure the assembled pipeline against a representative feed.** Corpus-median  $p_{50}$  and  $p_{99}$  of end-to-end latency, per operator-chosen tradeoff (Section 8.3).
7. **Flip one boundary at a time and re-measure.** Accept the flip if  $(p_{50}, p_{99})$  improves. Iterate until no single-boundary flip helps.
8. **Prior on where step 6 converges.** On Hawkes-clustered feeds, `async send` at every boundary whose removal would raise the largest stage; `fast_send` at every other boundary, in particular between sub-tasks downstream of the bottleneck whose combined service stays at or below  $s_{\max}$ , and between short sub-tasks below the floor. A hop that does not lower  $s_{\max}$  costs  $h$  at every quantile and compresses nothing.
9. **Do not pre-spin or pre-warm by default.** If pre-spinning or pre-warming is used, measure the ROI on a Hawkes-representative tape, net of the steady core-burn and cache-pollution costs. If the measurement improves the tail materially, keep the intervention. The paper’s prior is that on Hawkes feeds it will not (Section 9).

The workflow prioritises measurement over local reasoning at every step, on the specific grounds that the arrival law of the target feed is a first-order determinant of the optimum and is not a property of the code that can be introspected from the sender's call site.

## Part IV

# Conclusions

## 11 Conclusion

The paper’s two headline conclusions overturn the prevailing view in HFT system design. That view holds that the entire hot path should live on a single-thread event loop, on the grounds that every thread hop adds real per-message latency; the community is overwhelmingly in the single-thread camp. We show on real CME MDP3 packet flow that both halves of the argument fail.

**Conclusion 1: the thread-hop cost is a wash under real arrivals.** The dominant term in end-to-end latency on a Hawkes-clustered feed is not the mailbox handoff — it is the cluster-wait tail. On 3512 windows at  $h = 1.7 \mu\text{s}$  the corpus-median single-thread  $p_{99}$  is 2.2, 4.6, 12.8 and 41.5 times the service floor at  $T = 16, 32, 64, 128 \mu\text{s}$ , i.e. a tail excess of 11 to 2900 times the hop cost; the  $p_{99.9}$  excess is roughly twice that again. Adding a hop is  $+1.7 \mu\text{s}$  at every quantile; the tail it lets us compress is one to three orders of magnitude larger whenever  $T$  sits above the feed’s tight-gap floor. Below that floor the statement reverses and the hop is pure cost, which is why the rule this paper recommends is a threshold rule and not a blanket endorsement of splitting. The single-thread objection — “every hop is a fixed tax that no benefit can overcome” — is empirically wrong on this class of feed.

**Conclusion 2: two threads beat one on the tail at the cost of one hop on the median.** Splitting a service task into two half-length actors on separate threads joined by one  $1.7 \mu\text{s}$  mailbox hop lowers  $p_{99}$  on 100% of the 3512 windows at every service floor from 16 to  $128 \mu\text{s}$ , cutting the corpus-median tail excess to between 0.03 and 0.16 of its single-thread value for a median cost of exactly  $h$ ; the trade favours the split for any tail weight  $\kappa \gtrsim 0.095$  at  $T = 16 \mu\text{s}$  and  $\kappa \gtrsim 0.0004$  at  $T = 128 \mu\text{s}$ . At  $T \leq 8 \mu\text{s}$ , at or below the feed’s tight-gap floor, the split loses on  $p_{99}$ , as the design equation says it must. The pattern extends to  $N > 2$ : an  $N$ -stage tandem contracts the tail excess by a factor of  $N$  or better on every sample path (Theorem 1) — and in practice by far more than  $N$ , because a split that drives the per-stage service under the tight-gap floor removes the excess outright rather than dividing it — while paying  $(N - 1)h$  at every quantile. The  $p_{99}$ -minimising stage count rises with the task length, from 1 at  $T \leq 8 \mu\text{s}$  to 2 at 16, 4 at 32 and 8 at 64 and  $128 \mu\text{s}$ , up to the number of natural cut points in the decode chain (four or five on CME). The single-thread event loop is dominated on the tail on Hawkes-clustered feeds above the floor, and holds its advantage only on the median, by one hop per stage.

**Why the single-thread rule persists.** We think the field arrived at the single-thread rule under a mental model of Poisson-like arrivals, where clustering does not exist and, at the service times a hot path runs at, the tail sits on the service floor. On such a feed the rule is correct — our own results give exactly that answer under the Poisson null, for every task up to  $16 \mu\text{s}$ . What the field missed is that real exchange feeds have never been Poisson: near-critical branching ratios on CME futures ( $n \approx 0.9$  on ES trades in the literature,  $n \approx 0.8$  on NQ packets in this corpus) have been documented for over a decade, and it is precisely those clusters that reverse the architectural recommendation. When your feed is bursty, thread-hop overhead is the price of admission to an operating point with a far shorter tail at a bounded median cost.

**What this changes in practice.** Multi-actor pipelines are the correct architecture for HFT feeds specifically because those feeds are self-exciting. Framework designers should measure busy-poll dequeue on pinned isolated cores against the wait/notify mailbox on a Hawkes-representative tape rather than assume it wins (Section 9), drive the mailbox mutex uncontended, keep memory-pool discipline, and cut the decode pipeline at whichever of its natural boundaries (4–5 on CME) minimise the largest stage, with that stage at ingress and everything downstream merged up to its service time; the tail factor is  $s_{\max}/T$ , and a hop that does not lower  $s_{\max}$  is pure cost. Strategy designers of both POV-style and race-sensitive algorithms benefit from the same  $N$  budget for the same reason: the tail is what matters, and the pipeline compresses it.

**Retraction of a prior design context, and the corrected design rule.** An earlier paper by the present author [32] did not itself argue for the single-thread servicing chain; it took the collapsed single-thread servicing chain as the widely accepted design context of the time and offered `fast_send` as the mechanism that makes actor composition inside it free of context-switch overhead. What the present paper retracts is the design context, not the mechanism: on Hawkes-clustered feeds the single-thread servicing chain should no longer be the unconditional default. The corrected position is that `fast_send` (synchronous, inline, sub-microsecond) and asynchronous `BQueue` handoffs (a real pipelined hop at a few microseconds) are *complementary* tools, and the choice between them is a per-boundary design decision that depends on the service time being chained and the arrival law of the feed the pipeline sits on. Designers should use `fast_send` and asynchronous sends judiciously, according to which is appropriate for the task at hand.

In shipping KASPAR-HFT the `fast_send` vs asynchronous-send choice at each actor boundary is normally left open until the end of the build, and is resolved by measuring the end-to-end median and tail latency of the assembled pipeline under the target feed. A convenient property of the KASPAR-HFT actor API is that the receiving actor’s `MESSAGE_HANDLER` is agnostic to how it was invoked: swapping a synchronous `fast_send` at a call site for an asynchronous `send` (or vice versa) is a one-line change on the sender side and requires no code change on the receiver side. The choice therefore lives entirely at each individual call site and can be flipped per boundary at end-of-build with minimal engineering cost. The paper’s design equation now supplies the arrival-law-aware default value of that end-of-build measurement: on Hawkes-clustered feeds, long tasks should be split into shorter actors joined by asynchronous sends as a matter of preference, on the presumption that splitting compresses the overall tail more than the added hop cost inflates the median. The rule presupposes that the task being split has a service time  $T$  comfortably above the floor of Section 4.3; a task of order  $h$  or below cannot be usefully split (a  $1\ \mu\text{s}$  task cannot be turned into two  $500\ \text{ns}$  actors when a single hop alone costs  $1.7\ \mu\text{s}$ ). Short synchronous chains inside one actor’s servicing of one message therefore continue to use `fast_send` unchanged. The end-of-build measurement remains the authoritative arbiter of any specific boundary.

**Main caveats.** The corpus is a single instrument (NQ front-month, chan 318); ES and BTC replications are follow-on work. Constant per-packet service is a simplification; Section 6.6 removes it for the one source of variability the live data lets us measure, the message count per datagram, and finds the design rule unchanged, but per-message work that varies with book state is still not modelled. The cache-pollution cost of spinning receivers on many cores is bounded but not measured. And Theorem 1 is a bound whose slack depends on the arrival law; the paper measures that slack ( $\gamma = 3.3$  and  $2.7$  where it is defined, against the guaranteed  $\gamma \geq 1$ ) rather than deriving it, and a closed form for the slack on Hawkes input is open.

None of these caveats reverse the two conclusions on the 276-session corpus.

## 12 Future work

- Multi-instrument replication (ES, BTC).
- Closed form for the slack in Theorem 1 under Hawkes input, i.e. the exponent  $\gamma - 1$  as a function of  $(\rho, n, \beta T)$ . By the reduction this is a single-server question about  $M^{[K]}/D/1$  with Borel batches at two service levels, which is tractable; the paper measures  $\gamma$  instead.
- Span/clustering coupling on a coalescing-heavy feed. The span-aware rerun of Section 6.6 finds the serialisation and queueing channels close to independent on NQ, matching the live NQ streams; the live ZN streams show span rising steeply with queue depth instead. Whether the two channels compound on an instrument of that kind, and what that does to the design rule, is not answered here.
- Cache-pollution empirical measurement per spinning-receiver stage count.
- Power-law-kernel Hawkes fit versus exponential.
- Overnight-session regime.
- Own-capture pcap validation of the network sub-stage.
- Load-conditioned hop cost  $h(\rho)$ . This paper calibrates  $h = 1.7 \mu\text{s}$  from the low-load wake-latency figure in Maciejewski [32]; under sustained-load conditions inside a Hawkes cluster the receiver never sleeps and the effective  $h$  collapses toward the mutex + cache-line-transfer floor of a few hundred nanoseconds (Section 4.7). Measuring  $h(\rho)$  directly and re-running the sweep against the load-conditioned  $h$  would sharpen the reported tail-win, which the current single- $h$  sim understates.

## 13 Reproducibility

All code referenced below lives in the KASPAR-HFT repository [33], <https://github.com/vincent212/kaspar-hft>, which also carries the actor framework the paper’s recommendations are stated for.

Code: `arrival_paper/qsim_prep.py` (stage 1: per-window packet-arrival extraction keyed on `sendingTime`, and the packet-Hawkes fit, cached to `metadata.parquet` plus per-window arrival arrays; `--window-min` selects the window length) and `arrival_paper/qsim_run.py` (stage 2: the  $(T, h, N)$  tandem Lindley grid with the Poisson null and the per-window Theorem 1 bound check, minutes per full-corpus sweep under `numba`); `arrival_paper/hawkes_smoke.py` holds the Hawkes MLE and `arrival_paper/make_tables.py` regenerates Tables 1 and 2 and the derived quantities quoted in the text directly from the grid output, `arrival_paper/make_figure.py` regenerates Figure 2 from the same grid, and `arrival_paper/make_diagnostics.py` recomputes the Section 3.2 Poisson-rejection diagnostics (gap fractions, Fano factors, Hurst exponents, and the gap-shuffled surrogate) from the cache. The span-aware rerun of Section 6.6 is `arrival_paper/make_spans.py` (per-packet span, aligned to the cached arrivals), `arrival_paper/qsim_span.py` (the three service models and the span-shuffle arm) and `arrival_paper/make_span_tables.py`. Every number and every figure in Part II is therefore reproducible by one command. We note this because it was not previously true: the tandem-curves figure had no generating script in the repository, and silently went stale against a rerun of the sweep — which is exactly the failure mode a reproducibility section exists to prevent.

The two-stage split is what makes the sweeps in this paper practical: the expensive work (tape parsing and 3512 Hawkes MLEs) is done once and cached, after which a new  $(T, h, N)$  grid costs minutes rather than hours. The refined threshold sweep of Section 4.5 and the five-minute rerun of Section 4.6 were both produced this way.

Determinism: the Poisson null’s per-window seed is derived from an MD5 of `(session, window_id)` rather than from Python’s built-in `hash`, which is randomised per interpreter run, so the cache is byte-reproducible across re-runs; the metadata frame is sorted by `(session, window_id)` before writing for the same reason.

The live measurement of Section 6 is separate instrumentation on the production system rather than part of this pipeline: `MsgRec / msg_record()` in `frame/perf/act/LatencyProbe.hpp` emits one 16-byte record per admitted message carrying  $(t_1, \text{latency}, \text{qlen}, \text{idx})$ , and the readers `kaspr/perf/kh_*.py` on the `perf/live-wire-to-book` branch produce the per-message, packet-size and arrival-process tables quoted there. Its latency results are published in Maciejewski [32].

Data: 281 per-session message tapes derived from CME MDP3 channel 318 pcaps (NQ front month, 2025-01 through 2026-02, holidays and half-days excluded); the tapes are derived from licensed exchange data and are not redistributed, but the extraction script and the cached per-window metadata are released with the paper.

## References

- [1] Hawkes, A. G. (1971). Spectra of some self-exciting and mutually exciting point processes. *Biometrika* 58, 83–90.
- [2] Hawkes, A. G., Oakes, D. (1974). A cluster process representation of a self-exciting process. *Journal of Applied Probability* 11(3), 493–503.
- [3] Tanner, J. C. (1961). A derivation of the Borel distribution. *Biometrika* 48(1–2), 222–224.
- [4] Asmussen, S. (2003). *Applied Probability and Queues*, 2nd ed. Springer. Chapter X (heavy-tailed and subexponential asymptotics for single-server queues).
- [5] Gao, X., Zhu, L. (2018). Functional central limit theorems for stationary Hawkes processes and application to infinite-server queues. *Queueing Systems* 90, 161–206.
- [6] Jackson, J. R. (1957). Networks of waiting lines. *Operations Research* 5, 518–521.
- [7] Burke, P. J. (1956). The output of a queuing system. *Operations Research* 4, 699–704.
- [8] Kelly, F. P. (1979). *Reversibility and Stochastic Networks*. Wiley.
- [9] Kleinrock, L. (1976). *Queueing Systems, Vol. II: Computer Applications*. Wiley.
- [10] Whitt, W. (1983). The queueing network analyzer. *Bell System Technical Journal* 62, 2779–2815.
- [11] Friedman, H. D. (1965). Reduction methods for tandem queueing systems. *Operations Research* 13(1), 121–131.
- [12] Lindley, D. V. (1952). The theory of queues with a single server. *Mathematical Proceedings of the Cambridge Philosophical Society* 48(2), 277–289.
- [13] Foss, S., Korshunov, D. (2000). Steady-state asymptotics for tandem, split-match and other feedforward queues with heavy-tailed service. *Queueing Systems* 35, 65–80.

- [14] Baccelli, F., Foss, S., Lelarge, M. (2005). Tail asymptotics for monotone-separable networks. *J. Appl. Probability* 42(3), 793–812.
- [15] Heindl, A. (2001). Decomposition of general tandem queueing networks with MMPP input. *Performance Evaluation* 44(1–4), 5–23.
- [16] Kim, Y. H., Chydzinski, A. (2002). Connection-wise end-to-end performance analysis of queueing networks with MMPP inputs. *Performance Evaluation* 47(2–3), 137–156.
- [17] Dębicki, K., Mandjes, M. (2015). *Queues and Lévy Fluctuation Theory*. Springer.
- [18] Daw, A., Pender, J. (2018). Queues driven by Hawkes processes. *Stochastic Systems* 8(3), 192–229.
- [19] Koops, D. T., Boxma, O. J., Mandjes, M. R. H. (2018). Infinite-server queues with Hawkes arrivals. *Queueing Systems* 92, 27–52.
- [20] Chen, X., Blanchet, J. (2021). Perfect sampling of Hawkes processes and queues with Hawkes arrivals. *Stochastic Systems* 11(3), 264–283.
- [21] Sheldon, D., et al. (2023). Steady-state analysis and online learning for queues with Hawkes arrivals. arXiv 2311.02577.
- [22] Gao, X., Zhu, L. (2024). Single-server queues with state-dependent Hawkes arrivals. *Mathematics of Operations Research*.
- [23] Ozaki, T. (1979). Maximum likelihood estimation of Hawkes’ self-exciting point processes. *Ann. Inst. Statist. Math.* 31(1), 145–155.
- [24] Filimonov, V., Sornette, D. (2012). Quantifying reflexivity in financial markets: toward a prediction of flash crashes. *Physical Review E* 85, 056108.
- [25] Hardiman, S. J., Bercot, N., Bouchaud, J.-P. (2013). Critical reflexivity in financial markets: a Hawkes process analysis. arXiv 1302.1405.
- [26] Bacry, E., Mastromatteo, I., Muzy, J.-F. (2015). Hawkes processes in finance. *Market Microstructure and Liquidity* 1(1).
- [27] Byrd, D., Hybinette, M., Balch, T. H. (2020). ABIDES: Towards high-fidelity market simulation for AI research.
- [28] Frey, S., Li, K., Nagy, P., Sapora, S., Lu, C., Zohren, S., Foerster, J., Calinescu, A. (2023). JAX-LOB. arXiv 2308.13289.
- [29] Aquilina, M., Budish, E., O’Neill, P. (2022). Quantifying the high-frequency trading “arms race”. *Quarterly Journal of Economics* 137(1), 493–564.
- [30] Dean, J., Barroso, L. A. (2013). The tail at scale. *Communications of the ACM* 56(2), 74–80.
- [31] Zhao, K., Alizadeh, M., Pan, R., Ports, D. R. K., Perry, J., Yang, X., et al. (2023). Parsimon: scalable tail latency estimation for data center networks. NSDI 2023.
- [32] Maciejewski, V. (2026). Adapting the actor model of concurrency for high-frequency trading: synchronous message delivery (`fast_send`) and a tick-to-book latency study. arXiv:2609.21173. <https://arxiv.org/abs/2609.21173>

- [33] Maciejewski, V. (2026). KASPAR-HFT: an open-source C++20 actor framework and trading system for CME futures. Source repository. <https://github.com/vincent212/kaspar-hft>
- [34] Thompson, M. (2011). The Single Writer Principle. *Mechanical Sympathy*, 29 September 2011. <https://mechanical-sympathy.blogspot.com/2011/09/single-writer-principle.html>
- [35] Fowler, M. (2011). The LMAX Architecture. 12 July 2011. <https://martinfowler.com/articles/lmax.html>
- [36] LMAX Exchange (2011). Disruptor: high-performance alternative to bounded queues for exchanging data between concurrent threads. <https://lmax-exchange.github.io/disruptor/>
- [37] Rigtorp, E. (2020). Low latency tuning guide. <https://rigtorp.se/low-latency-guide/>
- [38] Chronicle Software (2023). Building fast trading engines: Chronicle’s approach to low-latency trading. <https://chronicle.software/insights/blogs/building-fast-trading-engines-chronicles-approach-to-low-latency-trading>
- [39] Real Logic (2016). Aeron: efficient reliable UDP unicast, UDP multicast, and IPC message transport. <https://aeron.io/>
- [40] Kwan, A., Wang, H., Luk, W. (2023). C++ design patterns for low-latency applications including high-frequency trading. arXiv 2309.04259.
- [41] Hu, W. (2024). The deterministic event-driven sequencer architecture: a competitive edge for high-frequency trading. Medium, October 2024. <https://medium.com/@hu.wenzhe124124/the-deterministic-event-driven-sequencer-architecture-a-competitive-edge-for-high-frequency>
- [42] Schmidt, D. C. (1995). Reactor: an object behavioral pattern for concurrent event demultiplexing and event handler dispatching. In *Pattern Languages of Program Design*, J. O. Coplien and D. C. Schmidt (eds.), Addison-Wesley.
- [43] Kegel, D. (1999). The C10K problem. <https://www.kegel.com/c10k.html>
- [44] Sysoev, I. (2004). nginx: a high-performance HTTP server and reverse proxy. <https://nginx.org/>
- [45] Dahl, R. (2009). Node.js: JavaScript runtime built on the V8 engine. Presented at JSConf EU, Berlin, November 2009. <https://nodejs.org/>
- [46] Welsh, M., Culler, D., Brewer, E. (2001). SEDA: an architecture for well-conditioned, scalable internet services. *Proc. 18th ACM Symposium on Operating Systems Principles (SOSP)*, 230–243.
- [47] Welsh, M. (2010). A retrospective on SEDA. <https://matt-welsh.blogspot.com/2010/07/retrospective-on-seda.html>
- [48] Kohler, E., Morris, R., Chen, B., Jannotti, J., Kaashoek, M. F. (2000). The Click modular router. *ACM Transactions on Computer Systems* 18(3), 263–297.
- [49] Murray, D. G., McSherry, F., Isaacs, R., Isard, M., Barham, P., Abadi, M. (2013). Naiad: a timely dataflow system. *Proc. 24th ACM Symposium on Operating Systems Principles (SOSP)*, 439–455.

- [50] Marty, M., de Kruijf, M., Adriaens, J., Alfeld, C., et al. (2019). Snap: a microkernel approach to host networking. *Proc. 27th ACM Symposium on Operating Systems Principles (SOSP)*, 399–413.
- [51] Kaffes, K., Chong, T., Humphries, J. T., Belay, A., Mazières, D., Kozyrakis, C. (2019). Shinjuku: preemptive scheduling for  $\mu$ second-scale tail latency. *Proc. 16th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 345–359.
- [52] Ousterhout, A., Fried, J., Behrens, J., Belay, A., Balakrishnan, H. (2019). Shenango: achieving high CPU efficiency for latency-sensitive datacenter workloads. *Proc. 16th USENIX NSDI*, 361–378.
- [53] Fried, J., Ruan, Z., Ousterhout, A., Belay, A. (2020). Caladan: mitigating interference at microsecond timescales. *Proc. 14th USENIX OSDI*, 281–297.
- [54] Chen, H., Delimitrou, C., Kozyrakis, C. (2021). Perséphone: a microsecond-scale time-sharing scheduler for wide service-time distributions. *Proc. 28th ACM SOSP*, 653–668.
- [55] DPDK Project (2024). DPDK Programmer’s Guide: run-to-completion and pipeline modes. [https://doc.dpdk.org/guides/prog\\_guide/](https://doc.dpdk.org/guides/prog_guide/)
- [56] Zhang, X. (2016). Pipelined FAST protocol market-data decoder for high-frequency trading. Chinese patent CN105654383A.
- [57] Vinnikov, A. (2014). Batch latency analysis and phase transitions for a tandem of queues. arXiv 1403.2400.
- [58] David, T., Guerraoui, R., Trigonakis, V. (2013). Everything you always wanted to know about synchronization but were afraid to ask. *Proc. 24th ACM Symposium on Operating Systems Principles (SOSP)*, 33–48.
- [59] Belay, A., Prekas, G., Klimovic, A., Grossman, S., Kozyrakis, C., Bugnion, E. (2014). IX: a protected dataplane operating system for high throughput and low latency. *Proc. 11th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 49–65.

# Appendices

## A The Borel batch-size tail at fixed branching ratio

This appendix records the calculation behind Equation (8) and the numerical check that the critical  $k^{-1/2}$  survival tail of the Borel distribution is confined to a finite window at every  $n < 1$ .

Stirling's formula  $k! = \sqrt{2\pi k} k^k e^{-k} (1 + O(1/k))$  in Equation (7) gives

$$\Pr(K = k) = \frac{(nk)^{k-1} e^{-nk}}{k!} = \frac{n^{k-1} k^{k-1} e^{-nk}}{\sqrt{2\pi k} k^k e^{-k}} (1 + O(1/k)) = \frac{1}{n\sqrt{2\pi}} k^{-3/2} (ne^{1-n})^k (1 + O(1/k)).$$

Writing  $n = 1 - \epsilon$ ,  $\log(ne^{1-n}) = \log(1 - \epsilon) + \epsilon = -\epsilon^2/2 - \epsilon^3/3 - \dots$ , so the geometric factor is  $\exp(-k\epsilon^2/2 (1 + O(\epsilon)))$ . The pmf is therefore  $\propto k^{-3/2}$  for  $k \ll k^* = \epsilon^{-2}$  and geometric beyond; summing,  $\Pr(K > k) \asymp k^{-1/2}$  on the same window and  $\asymp k^{-3/2} e^{-k\epsilon^2/2}$  beyond it. At  $n = 1$  exactly,  $k^* = \infty$ ,  $\Pr(K > k) \sim \sqrt{2/\pi} k^{-1/2}$  (Borel–Tanner), and  $E[K] = \infty$ ; that is the unstable case  $\rho = \infty$ , which no stationary queue attains.

Because the tail is geometric,  $K$  has finite exponential moments  $E[e^{\theta K}] < \infty$  for  $\theta < \epsilon^2/2 (1 + O(\epsilon))$  and is therefore outside the subexponential class  $\mathcal{S}$ . Heavy-tail queueing results that assume subexponential batch work, including the single-big-batch principle for  $M^{[K]}/G/1$  and the monotone-separable tail theorem of Baccelli et al. [14], do not apply, and the stationary waiting time of the batch-arrival queue has a Cramér–Lundberg (exponential) tail at every fixed  $n < 1$ .

Table 10 reports the local log-log slope of the exact pmf,  $\Delta \log \Pr(K = k) / \Delta \log k$  between successive  $k$ , which a pure  $k^{-3/2}$  law would hold at  $-1.5$ .

| $n$  | $k^*$ | $10 \rightarrow 30$ | $30 \rightarrow 100$ | $100 \rightarrow 300$ | $300 \rightarrow 1000$ | $1000 \rightarrow 3000$ |
|------|-------|---------------------|----------------------|-----------------------|------------------------|-------------------------|
| 0.80 | 25    | −1.92               | −2.84                | −5.71                 | −14.96                 | −43.63                  |
| 0.90 | 100   | −1.59               | −1.81                | −2.48                 | −4.62                  | −11.26                  |
| 0.95 | 400   | −1.52               | −1.57                | −1.73                 | −2.25                  | −3.85                   |
| 0.99 | 10000 | −1.50               | −1.50                | −1.51                 | −1.53                  | −1.59                   |

**Table 10.** Local log-log slope of the Borel pmf, Equation (7), between successive cluster sizes  $k$ . The critical slope  $-1.5$  holds while  $k \ll k^* = (1 - n)^{-2}$  and steepens geometrically beyond.

The practical reading is that at the branching ratios fitted on the corpus ( $n_{\text{pkt}} \approx 0.8$ , Section 3.2),  $k^* \approx 25$ : clusters of order a hundred packets are the largest that occur with non-negligible frequency, and the single-stage wait behind the largest clusters has scale  $Tk^* \approx 25T$ , against the corpus-median  $p_{99}$  tail excess of Section 4, which is zero below the tight-gap floor and rises from  $1.2T$  at  $T = 16 \mu\text{s}$  to  $39T$  at  $T = 128 \mu\text{s}$ . This is the object that Theorem 1 divides by  $N$ , and the reason the  $n = 0.80$  row of Table 10 is the relevant one.

## B Numerical check of Theorem 1

A direct simulation of the model of Theorem 1 checks both the identity Equation (4) and the bound Equation (6). Immigrants arrive as a Poisson process at rate  $\mu = 0.004 \mu\text{s}^{-1}$ , each spawning a Galton–Watson cluster with Poisson(0.9) offspring, so  $K \sim \text{Borel}(0.9)$  and  $E[K] \approx 10$ ; clusters are placed as instantaneous batches, giving  $\rho = \bar{\lambda}T \approx 0.40$  at  $T = 10 \mu\text{s}$ . The tandem has  $N = 4$  equal stages of  $2.5 \mu\text{s}$  and  $h = 1.7 \mu\text{s}$ ;  $2 \times 10^5$  messages are simulated with an explicit per-stage FIFO recursion

(independent of the paper’s `qsim_run.py`). Over all messages,  $\max_j |W_j^{(4)} - (w_j(T/4) + 3h)| < 10^{-9} \mu s$  and Equation (6) is violated by no message. The quantiles are

| $q$   | $W_q^{(1)}$ | $W_q^{(4)}$ | $W_q^{(1)}/4 + 3h$ | $(W_q^{(4)} - 3h)/W_q^{(1)}$ |
|-------|-------------|-------------|--------------------|------------------------------|
| 0.5   | 350.0       | 55.1        | 92.6               | 0.14                         |
| 0.9   | 2066.6      | 355.1       | 521.8              | 0.17                         |
| 0.99  | 5342.1      | 1127.6      | 1340.6             | 0.21                         |
| 0.999 | 7586.2      | 1585.1      | 1901.7             | 0.21                         |

in microseconds. The tandem lands below the bound at every quantile, with the ratio in the last column between 0.14 and 0.21 against the bound’s 0.25; the slack is the  $\rho/N$  utilisation effect described in the remark after Theorem 1. Under the Poisson null at the same message rate ( $\rho = 0.40$ , which is far above the corpus utilisation and is used here only to make the Poisson wait visible) the single-stage wait quantiles at  $q = 0.9, 0.99, 0.999$  are  $10.3, 24.8, 38.5 \mu s$  and the tandem’s are  $5.1, 7.5, 9.2 \mu s$ , again inside the bound. At the corpus utilisation ( $\rho = 0.0045$  at this appendix’s  $T = 10 \mu s$ ) the same two collapse to 0 and  $3h$  at  $q = 0.9$  and  $q = 0.99$ , which is Corollary 2; at  $q = 0.999$  they do not, since  $\rho$  there exceeds  $1 - q = 0.001$  and the corollary does not apply.

## C Numerical check of the unequal-stage claims

The same generator as Appendix B (Borel(0.9) clusters,  $\rho = 0.40$ ,  $T = 10 \mu s$ ,  $h = 1.7 \mu s$ ), with each cluster’s members spread uniformly over  $3 \mu s$  instead of placed as an instantaneous batch and  $4 \times 10^5$  messages, run through the partitions of Remark 2.2 with an explicit per-stage FIFO recursion. The last two columns are the tail-excess ratio  $(\Delta_{99}(P) - \text{hops} \cdot h)/\Delta_{99}(\text{single})$  and the bound  $s_{\max}/T$  from Theorem 1.

| partition      | $s_{\max}$ | hops | $p_{50}$ | $p_{99}$ | $p_{99.9}$ | ratio | bound |
|----------------|------------|------|----------|----------|------------|-------|-------|
| 10             | 10.0       | 0    | 388.2    | 5719.0   | 8249.3     | 1.00  | 1.00  |
| 7 + 3          | 7.0        | 1    | 213.6    | 3429.5   | 5297.1     | 0.60  | 0.70  |
| 3 + 7          | 7.0        | 1    | 213.6    | 3429.5   | 5297.1     | 0.60  | 0.70  |
| 5 + 5          | 5.0        | 1    | 134.9    | 2265.5   | 3574.2     | 0.39  | 0.50  |
| 7 + 2 + 1      | 7.0        | 2    | 215.3    | 3431.2   | 5298.8     | 0.60  | 0.70  |
| 4 + 3 + 3      | 4.0        | 2    | 104.6    | 1746.5   | 2714.6     | 0.30  | 0.40  |
| 4 + 3 + 2 + 1  | 4.0        | 3    | 106.3    | 1748.2   | 2716.3     | 0.30  | 0.40  |
| $2.5 \times 4$ | 2.5        | 3    | 66.4     | 1049.8   | 1562.2     | 0.18  | 0.25  |

Latencies are end-to-end in microseconds. The three claims of Remark 2.2 read off directly: the tail excess depends on the partition only through  $s_{\max}$  and sits below the bound at every row;  $7 + 3$  and  $3 + 7$  are identical to the reported precision; and  $7 + 2 + 1$  differs from  $7 + 3$ , and  $4 + 3 + 2 + 1$  from  $4 + 3 + 3$ , by exactly one hop cost  $h = 1.7 \mu s$  at every quantile. The  $3 \mu s$  intra-cluster spread, which is above  $s_{\max}$  for the last row, is why the equal-stage ratio here (0.18) sits slightly below the batch-input value of Appendix B (0.21): a cluster that arrives spread out lets the faster tandem drain part of it before the next member lands.